



TOHOKU
UNIVERSITY



Cyberscience
Center

High Performance Computing

高性能計算論

Volume 3

Cyberscience Center, Tohoku Univ

Hiroyuki Takizawa


<takizawa@tohoku.ac.jp>

Previous Lesson

■ Parallel Computers

- SIMD Overview
 - A single instruction operates on multiple data
 - Auto-vectorization to use vector units
- Shared-memory computers
 - Parallel computing with a single memory space
 - Communication via shared data
- Distributed-memory computers
 - Parallel computing with multiple memory spaces
 - Communication via explicit message passing
- Hierarchical (hybrid) systems
 - Mixture of shared-memory and distributed-memory parallel computers
 - Modern HPC systems
- Networks
 - No network topology is optimal in every regard

SIMD Overview



■ Vector Lane

- A pathway for each data item to be processed

■ Vector Width

- The width of the vector unit usually expressed in bits

■ Vector Length

- The number of data items being processed by a vector instruction

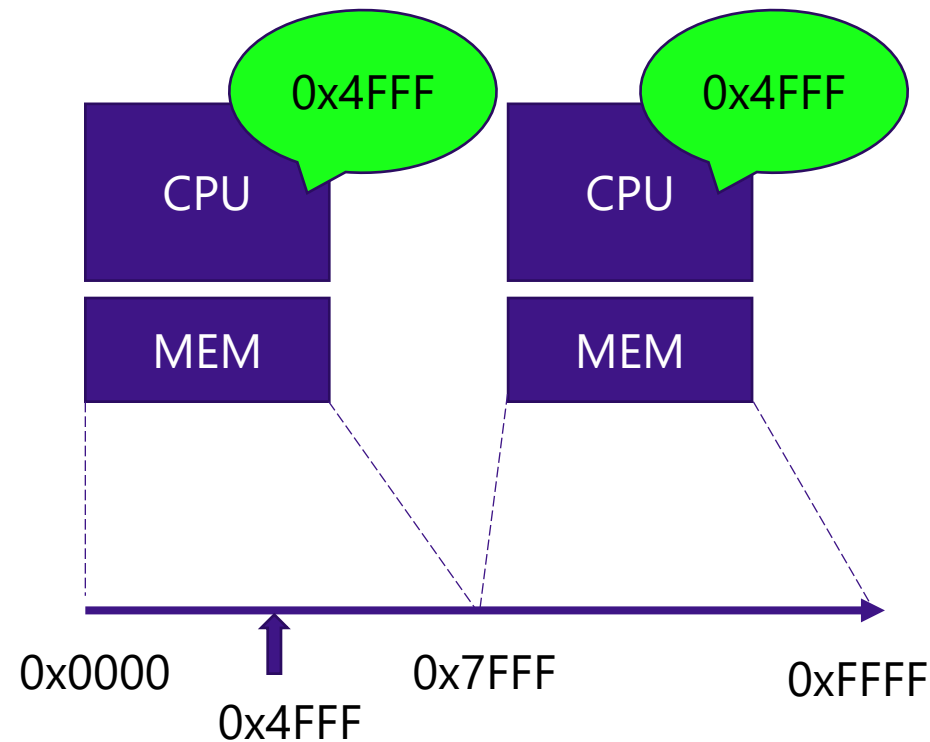
■ Vector (SIMD) instruction set

- The set of instructions to utilize the vector processing capability

Sharing a Memory Space

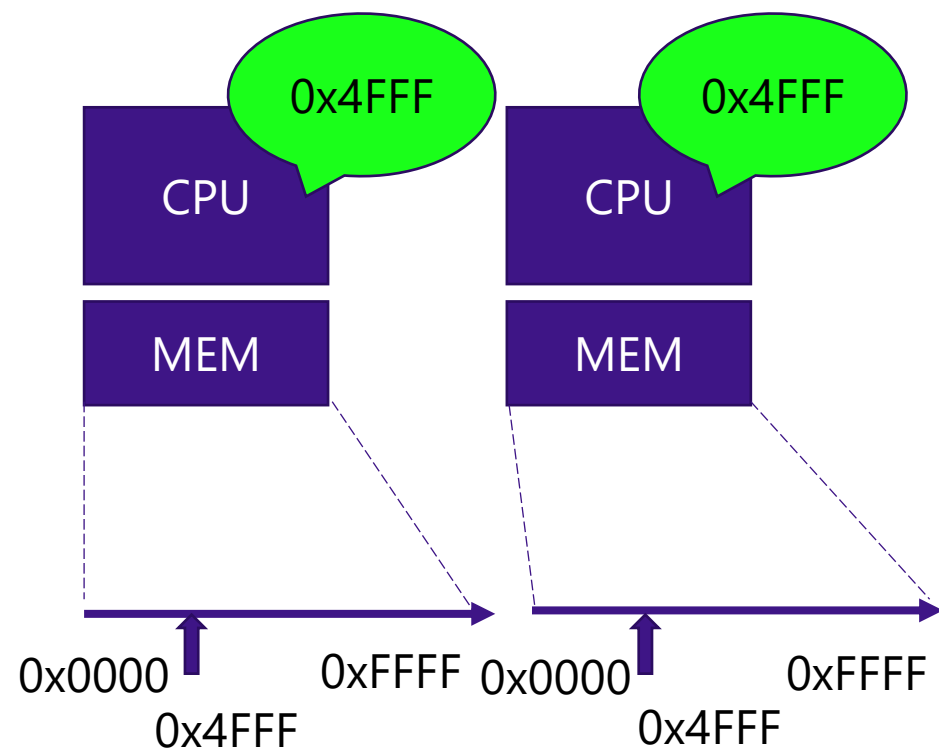
Shared-memory (NUMA)

Each memory device is mapped to a part of the memory space.




Distributed-memory

Each memory device has its own memory space.



Basic Network Topology

■ How nodes are connected via links.



Good point the bisection bandwidth scales with the network size, and a shorter latency for communication with neighbors.


Bad point the diameter is large and grows at $O(\sqrt{n})$ for n switches.

A torus network has a wraparound connection to reduce the diameter by half. In higher-dimensional networks a node can have direct connections with more nodes.

Basic Network Topology (cnt'd)


■ How nodes are connected via links.

Cut surface



Binary tree

Cut surface



Fat tree

Good point the diameter is small and grows slowly at $O(\log n)$.

Bad point the bisection bandwidth could be small

A fat tree network increases the bandwidth of upper-layer links to increase the bisection bandwidth → more hardware cost and adaptive routing mechanisms.

Parallel Algorithm Design

■ It's time to design a parallel algorithm.

- An **algorithm** is a well-defined procedure to solve a problem and consists of a finite number of instructions.
- A **parallel algorithm** is an algorithm designed for running on a parallel computer.
 - A parallel algorithm is implemented as a parallel program in a parallel programming language.
- **Parallelization** is to design an algorithm, program, or system to run in parallel.

There is no all-around way of parallelization.

Today's Topics

■ Job Level Parallelism

- What is Job?
- Job Scheduling

■ Parallel Algorithm Design

- A Model of Parallel Computation
- Foster's Design Methodology
- Case Studies

How to run a program on HPC?

■ Batch Job

- **A unit of work from user's point of view**
 - Submitted to an HPC System
- A job is usually a batch of tasks
 - Task is a unit of work for a computer

■ Front-end Server and Compute Nodes

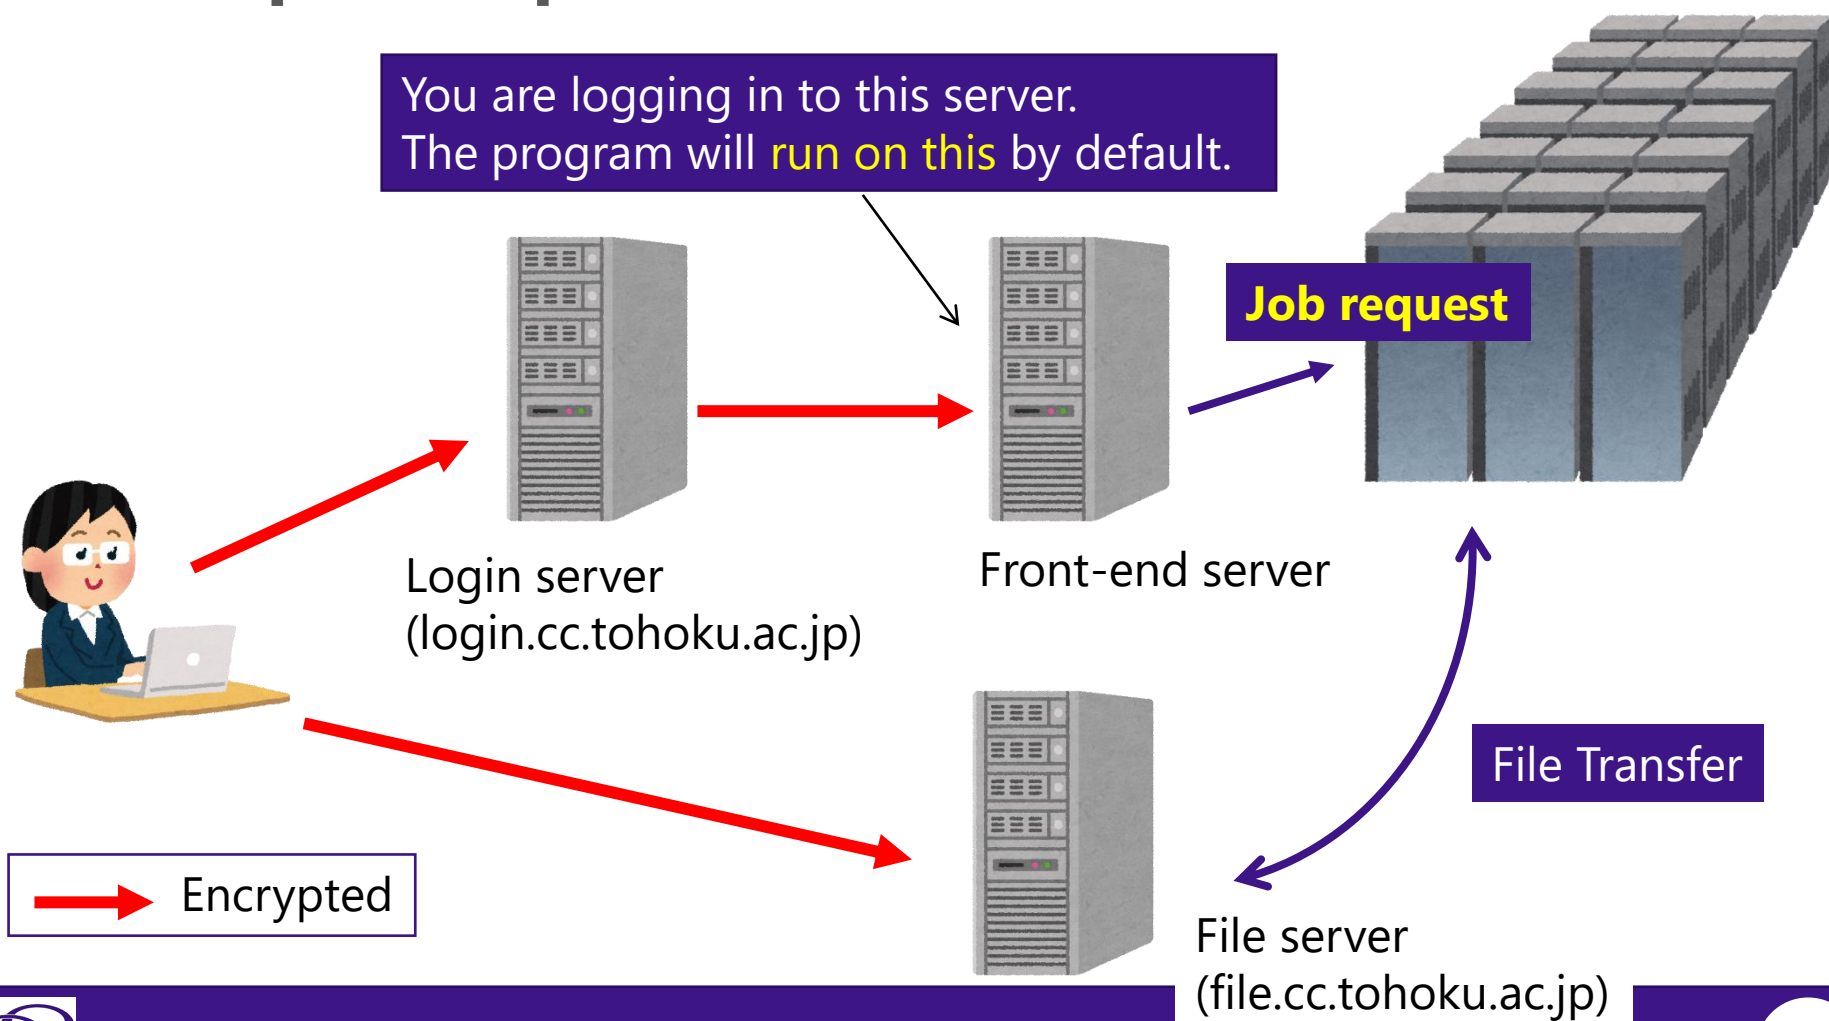
- Users can log-in to the front-end server, but not directly to the compute nodes
- How to run a job on compute nodes?
 - User submit a job from the front-end server for job execution.



Why Job Submission Needed?

■ Supercomputer AOBAs = **Shared** Resource

You are logging in to this server.
The program will **run on this** by default.



Job Script File

■ Script of expressing how to execute tasks

- A text file
- List of Linux commands to be executed

```
#!/bin/sh -  
#PBS -q lx_edu  
#PBS -l elapstim_req=0:05:00
```

This is a bourne shell job script.

The job is submitted to a queue named lx_edu.

Requesting job execution of 5 minutes

```
cd $PBS_O_WORKDIR
```

Job is executed at the directory it is submitted

```
# Linux commands  
echo "hello world"
```

The text following # is **basically** ignored at the execution. But some lines are used as hints for the job scheduler.

Job Submission


lx_edu is dedicated to this class.
Use **lx** usually.

■ Write a shell script file (run.sh)

```
#!/bin/sh -  
#PBS -q lx_edu  
#PBS -l elapstim_req=0:05:00  
cd $PBS_O_WORKDIR  
mpirun -np 4 ./sum1
```

This job will be executed on **AOBA-B**.

Waiting Queue named "lx_edu"
associated with AOBA-B



AOBA-B

AOBA-A

■ Submit it to job scheduler

```
qsub run.sh
```

Job submission

Your job will be appended at the end of the queue.

```
qstat
```

Check the status



■ Get the result

```
run.sh.eXXXXXX
```

stderr


```
run.sh.oXXXXXX
```

stdout

These files will be created in
the same directory.
xxxxx is the job ID. (5 digits)

Job Scheduling

- Decide **where** and **when** a job is executed
 - Necessary for efficient use of shared resources
 - The most basic policy is **First Come First Serve (FCFS)** policy.




Backfilling

■ Job can be allowed to overtake others if ...


- not delay the execution start of any other jobs.

= **Conservative Backfilling**



- not delay the execution start of the first waiting job.

= **EASY Backfilling**





Job 3 is delayed but the **makespan** becomes shorter.

The time length from the start of the first job to the end of the last job.

Workflows


■ Dependency between jobs/tasks



Spatial and temporal assignment of nodes must be considered to achieve the fastest possible execution under the dependency constraint.

→ NP hard problem and some heuristics are needed to get suboptimal ones.

Exploiting Parallelism




TOP500 List No.1 System

Difficult to scale more

■ Performance development slows down

PERFORMANCE DEVELOPMENT

TOP 500




Presented by Erich Strohmaier (June 28, 2021).




Importance of efficient use

■ Job Scheduling

- Efficient use of a limited amount of resource



System throughput improves by better job scheduling!



Technical challenges

■ Throughput Challenge

- Large ensemble simulations require massive numbers of jobs that cannot comfortably be ingested and scheduled by the traditional approach.

■ Co-scheduling Challenge

- Complex coupling requires sophisticated co-scheduling that the existing centralized approaches cannot easily provide.

■ Job coordination and communication challenge

- Intimate interactions with RJMS are required to keep track of the overall progress of the ensemble execution, and existing approaches lack well-defined interfaces.

■ Portability Challenge

- There has been a proliferation of ad hoc implementations of user-level schedulers as an attempt to tackle the above challenges. They are often non-portable and come with a myriad of side effects (e.g., millions of small files just to coordinate the current state of an ensemble).

Various Workloads

■ Rigid jobs (conventional)

- E.g. Numerical simulations ...

■ On-demand jobs

- E.g. Urgent jobs ...

■ Malleable jobs

- E.g. Big data analysis, parameter survey ...

■ Others

- Containerization
- Sensor data streaming

Is it possible to put them all together within a single system?

Today's Topics

■ Job Level Parallelism

- What is Job?
- Job Scheduling

■ Parallel Algorithm Design


- **A Model of Parallel Computation**
- Foster's Design Methodology
- Case Studies

Task/Channel Model

■ Task/Channel Model (Ian Foster, 1995)

- Parallel computation is a set of **tasks** sending messages via **channels**.
 - A task has its private data in the local memory.
 - sends local data values to other tasks via output ports.
 - receives data values from other tasks via input ports.

A channel is a message queue connecting one task's output port with another task's input port.



A task consists of a program, its local data, and a collection of I/O ports.

Synchronous / Asynchronous

■ What happens if a task tries to receive a value but no value is available?

- The task must wait until the value appears.
(= the task is **blocked**.)

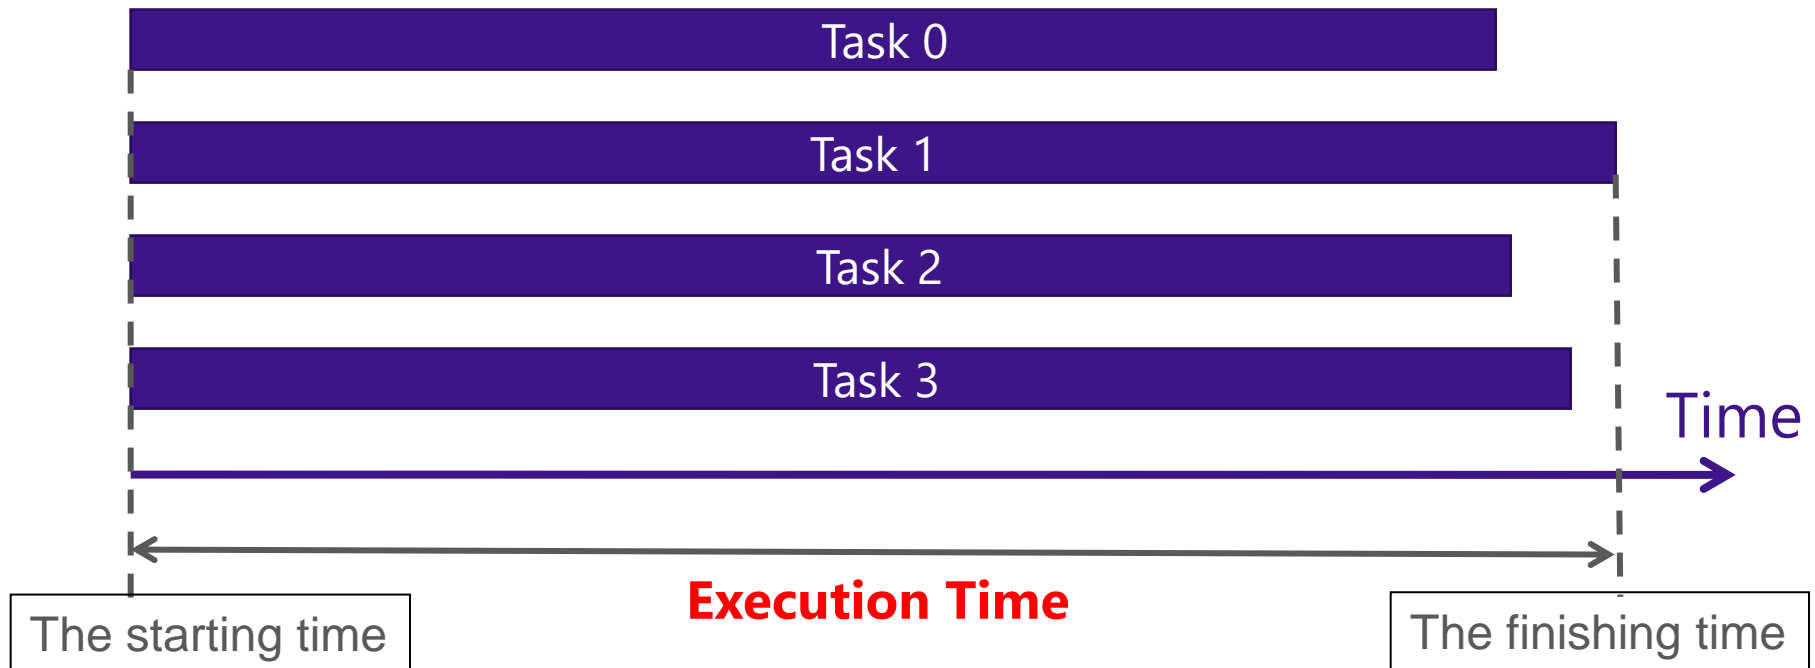
■ Synchronous/asynchronous communications

- Sending a message is never blocked in T/C model.
(= an **asynchronous** or **non-blocking** operation)
 - A task can send a message even if previous messages are not received yet.
- Receiving a message is a **synchronous** or **blocking** operation.

Execution Time

■ Execution time of a parallel algorithm:

- The period of time during **any task is active**.
(Shorter is Better.)



Today's Topics

■ Job Level Parallelism

- What is Job?
- Job Scheduling

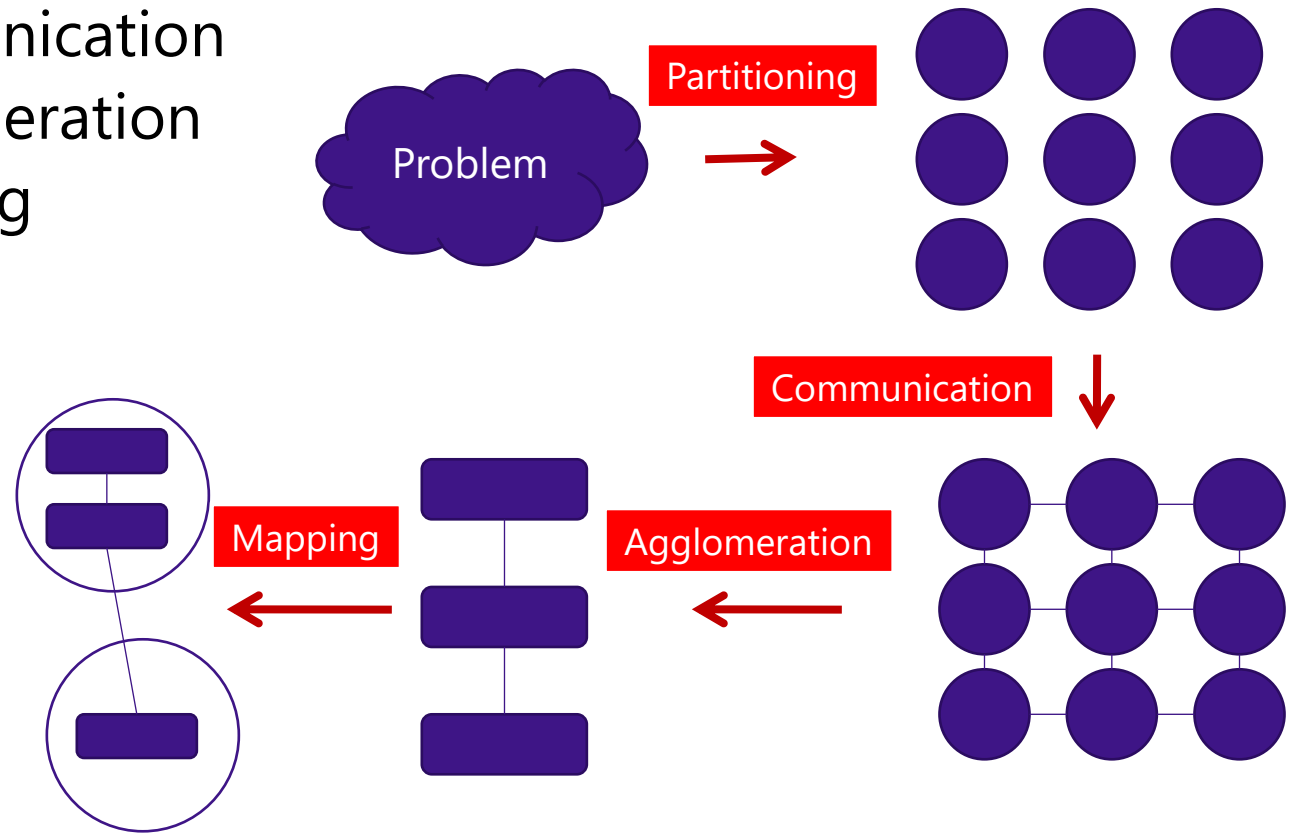
■ Parallel Algorithm Design

- A Model of Parallel Computation
- **Foster's Design Methodology**
- Case Studies

Foster's Design Methodology

■ Four steps for designing parallel algorithms

- Partitioning
- Communication
- Agglomeration
- Mapping



Partitioning

■ To discover as much parallelism as possible!

- Dividing the computation and the data into pieces.
= partitioning

■ Domain decomposition


- A data-centric approach
 - first divides the data into pieces, and then determines the computations with the data.

■ Functional decomposition

- A computation-centric approach
 - first divides the computation into pieces, and then determines how to associate data items with the individual computations.

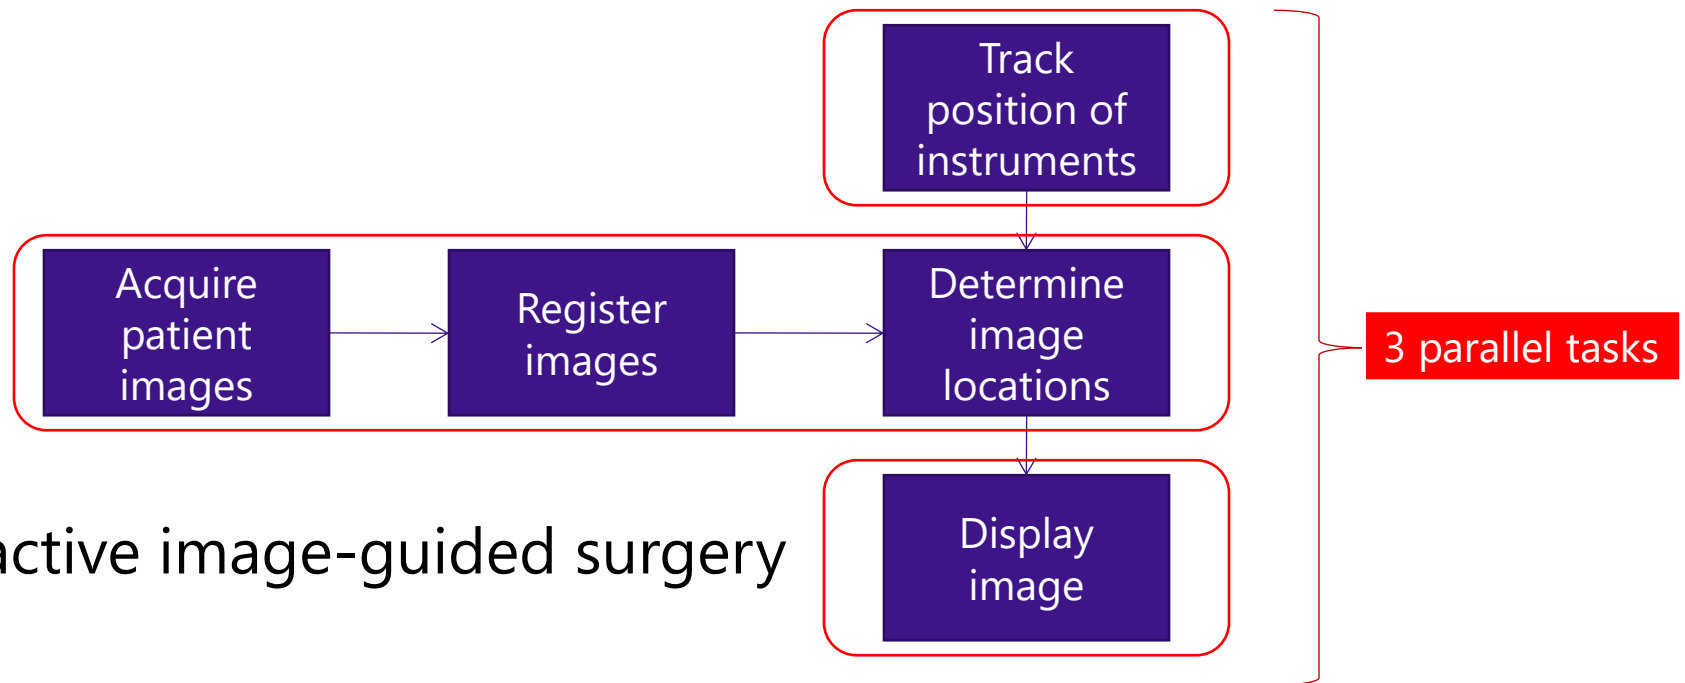
Domain Decomposition

- Suppose a 3D matrix is the largest and most frequently accessed.



Functional Decomposition

- **Whichever decomposition we choose, we call each of computation pieces a primitive task.**
 - The goal is to identify as many tasks as possible.



Interactive image-guided surgery

Partitioning Checklist

■ The best designs satisfy all of the following attributes:

- There are much more primitive tasks than processors in the target system (at least an order of magnitude).
 - If not, later design options may be too constrained.
- Redundant computation and redundant data structure storage are minimized.
 - If not, it may not work well when the problem size increases.
- Primitive tasks are roughly the same size.
 - If not, it may be hard to balance work among processors.
- The number of tasks is an increasing function of the problem size.
 - If not, it may be impossible to use more processors to solve larger problem instances.

Communication

■ **Communication:** determining the communication pattern among primitive tasks.

■ **Local and Global Communications**

- **Local communication:**
communication with a small number of other tasks.
- **Global communication:**
communication with a significant number of other tasks.
 - e.g. calculating the sum of values calculated by primitive tasks.
 - Communication channels for global communication is not drawn at this stage of the algorithm's design.

Communication Checklist

- **Communication = overhead of a parallel algorithm**
(not required by a sequential algorithm)
→ Minimizing parallel overhead is an important goal of parallel algorithm design.

- **Communication Quality Checklist**

- The communication operations are balanced among the tasks.
- Each task should communicate with only a small number of neighbors whenever possible.
- Tasks can perform their communications concurrently.
- Tasks can perform their computations concurrently.


Agglomeration

■ **Agglomeration**: grouping tasks into larger tasks


- considering a real parallel computer in mind.
- to improve performance or to simplify programming

■ **The goals of agglomeration**

- Lowering the communication overhead
 - **Increasing the locality**: agglomerating primitive tasks communicating with each other → reduction in the communication overhead



Eliminating communication




Reducing # msg transmissions



Agglomeration (Cont'd)

■ The goals of agglomeration (cont'd)

- Maintaining the scalability of the design
 - **Don't combine too many tasks**. A program should be portable to a system of more processors.



4 parallel tasks



2 parallel tasks

- Reducing the software engineering cost
 - allows us to make greater use of the existing sequential code.

Agglomeration Checklist

■ The best designs satisfy all of the following attributes.

- The agglomeration has increased the locality of the algorithm.
- Replicated computations take less time than the communications they replace.
- The amount of replicated data is small enough to allow the algorithm to scale.
- Agglomerated tasks have similar computational and communication costs.
- The number of tasks is an increasing function of the problem size.
- The number of tasks is as small as possible, yet as great as the number of processors in the target computers.
- The trade-off between the chosen agglomeration and the cost of modifications to existing sequential code is reasonable.

Mapping

■ **Mapping**: assigning tasks to processors.

- Automatic
 - MPI runtime automatically assigns tasks to a distributed-memory parallel computer.
 - OS automatically assigns tasks to shared-memory systems.
- Manual mapping sometimes improves performance
 - **MPI rank layout/mapping** on distributed-memory system
 - **Thread affinity** on shared-memory system


■ **The goals of mapping**

- Maximizing **processor utilization**.
 - Minimizing **interprocessor communication**.
- Those goals are **often conflicting**!

Processor Utilization

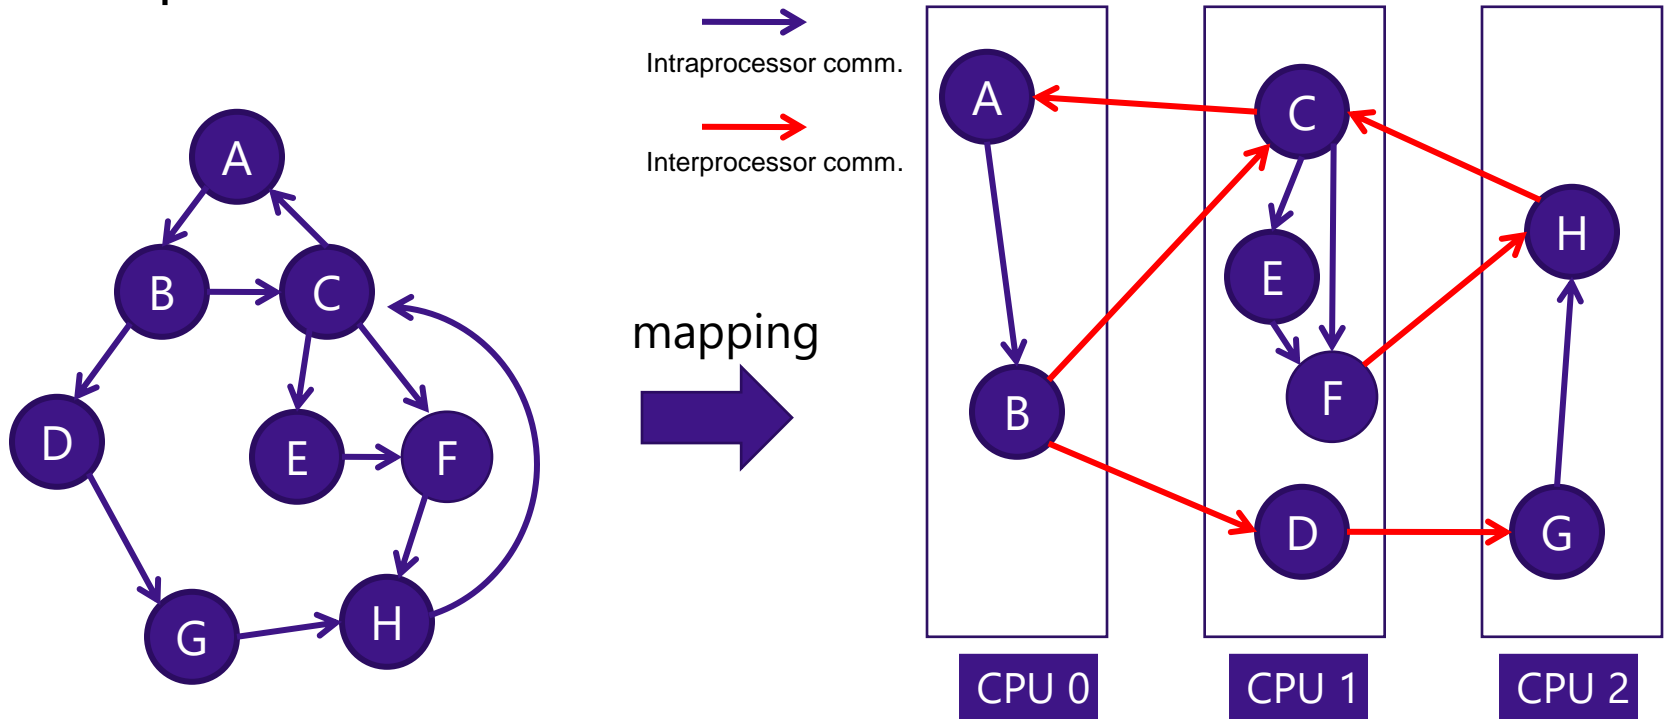
■ Average percentage of time **executing tasks necessary for solution** of the problem

- Is maximized when the computation is balanced evenly, i.e. all processors begin and finish the execution at the same time.
- decreases if one or more processors are idle.



Interprocessor Communication

- **increases when two tasks connected by a channel are mapped to different processors.**
 - decreases when they are mapped to the same processor.



How to Find a Good Mapping (1/2)

■ NP-hard problem = no polynomial-time algorithms
→ **we must rely on heuristics to find a reasonably**

good solution

- In the case of domain decomposition...
 - The tasks after agglomeration often have similar size
 - i.e. computational loads are balanced among tasks
 - A good strategy is to create p agglomerated tasks ($p = \#$ processors)
 - minimize the communication and map each of them to its own processors
- In the case of a fixed number of tasks...
 - Regular comm. and various task sizes → cyclic (or interleaved) mapping tasks to procs.
 - balancing the computational load at expense of higher communication costs
 - Unstructured comm. pattern → mapping tasks to minimize the comm. overhead


How to Find a Good Mapping (2/2)

■ Dynamic Load Balancing

- analyzes the current tasks and produces a new mapping of tasks to processors at runtime.
- needed when tasks are created and destroyed at runtime
- or, needed when the comm. or the comp. requirements vary widely


■ Task scheduling

- Suppose short-lived, independent tasks



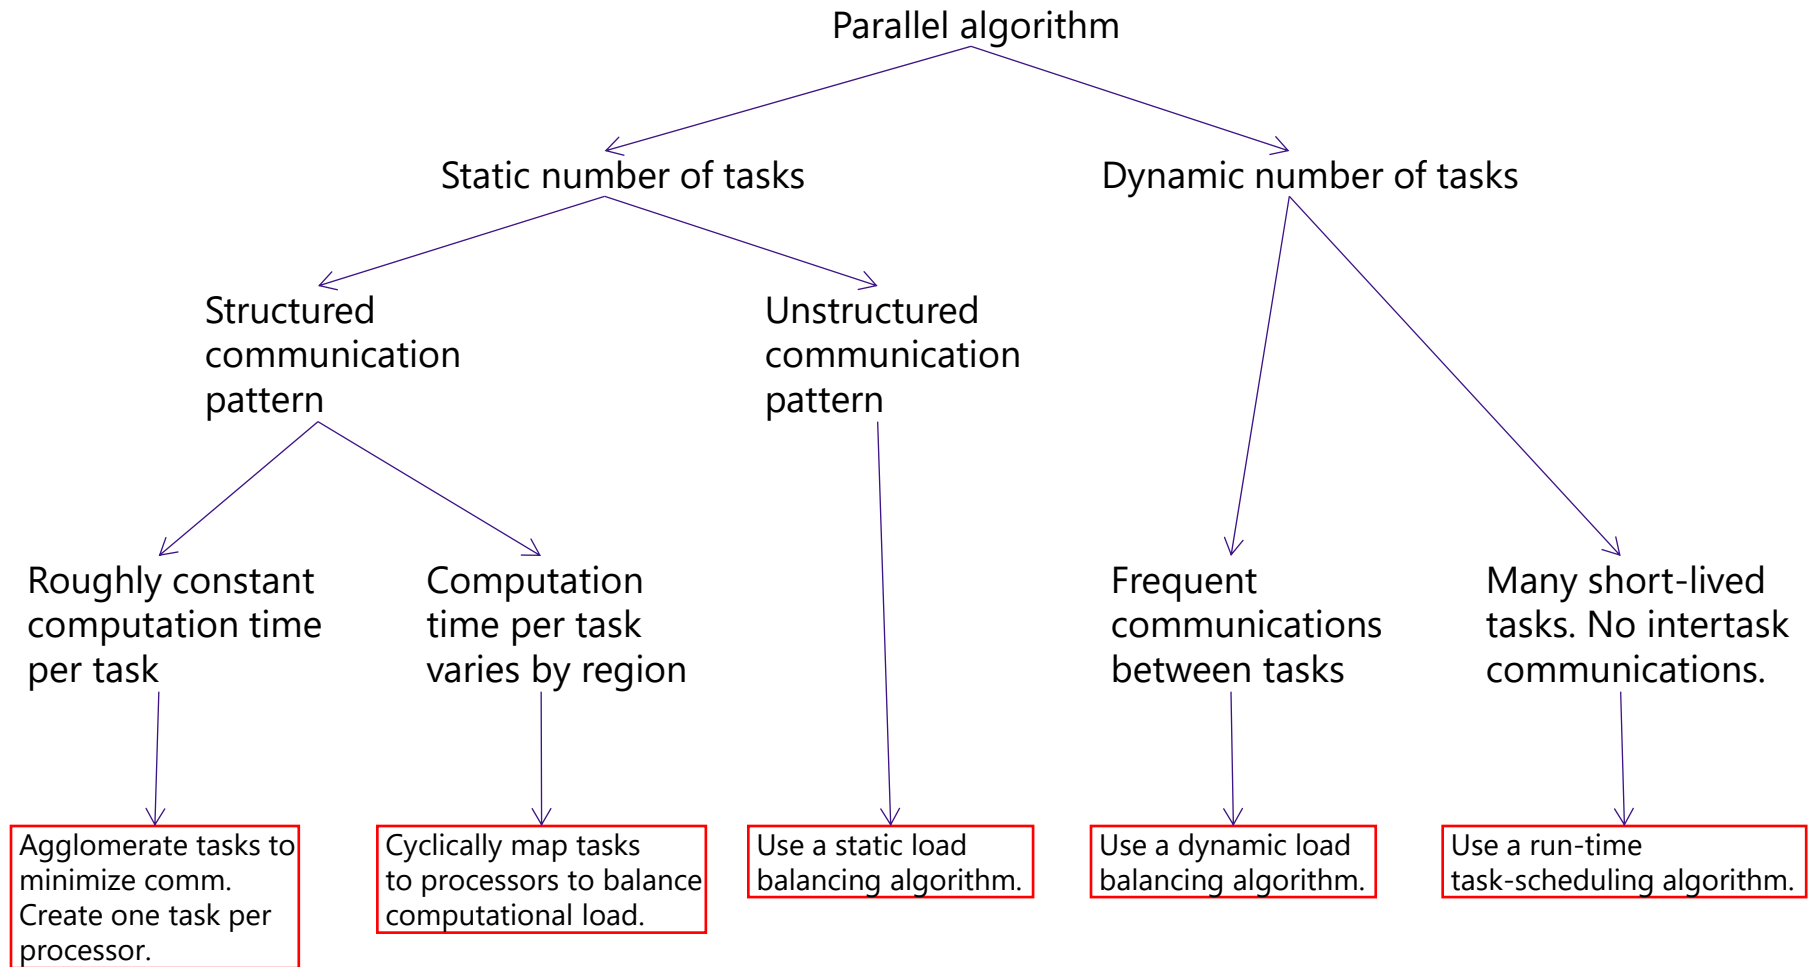
High Perf

Centralized Task Scheduling



Distributed Task Scheduling

Mapping Strategies



Mapping Checklist

■ The following checklist can help you decide if you've done a good job of design alternatives.

- Two designs, one task per processor and multiple tasks per processor, have been considered.
- Both static and dynamic allocations of tasks to processors have been evaluated.
- If a dynamic allocation of tasks to processors has been chosen, the manager (task allocator) is not a bottleneck to performance.
- If a static allocation of a tasks to processors has been chosen, the ratio of tasks to processors is at least 10:1.

Today's Topics



■ Job Level Parallelism

- What is Job?
- Job Scheduling

■ Parallel Algorithm Design

- A Model of Parallel Computation
- Foster's Design Methodology
- **Case Studies**

Boundary Value Problem



Difference Method (1/2)

■ Heat conduction

Temperature at location x and time t

$$\frac{\partial u(x, t)}{\partial t} = \frac{\partial^2 u(x, t)}{\partial x^2}$$

$$\frac{\partial u(x)}{\partial x} \approx \lim_{h \rightarrow 0} \frac{u(x + h) - u(x)}{h}$$

$$\frac{\partial^2 u(x)}{\partial x^2} \approx \lim_{h \rightarrow 0} \frac{u(x + h) - u(x)}{h^2} - \frac{u(x) - u(x - h)}{h^2}$$

$$= \lim_{h \rightarrow 0} \frac{u(x + h) - 2u(x) + u(x - h)}{h^2}$$

Difference Method (2/2)


■ Heat conduction

$$\frac{\partial u(x, t)}{\partial t} = \frac{\partial^2 u(x, t)}{\partial x^2}$$


$$\frac{u(t+k) - u(t)}{k} = \frac{u(x+h) - 2u(x) + u(x-h)}{h^2}$$

k and h are small numbers.

Temperature at time step $t+1$ can be calculated by using the current temperatures of each location and its neighbors.



Finite Difference Approximation



Partitioning

■ One data item per grid point


- The rod at each time step is decomposed into pieces (**grid points**) in a uniform way.

$u(0,t)$	$u(1,t)$...	$u(x,t)$...	$u(N-1,t)$
----------	----------	-----	----------	-----	------------

The rod is decomposed into N pieces.


■ Associate one primitive task with each grid point and each time step

- Two-dimensional domain decomposition
 - Time and location




Communication

- Identify communication pattern between primitive tasks
- Each interior primitive task has three incoming and three outgoing channels




Agglomeration and Mapping



(a)



(b)



(c)

Agglomeration



Finding the Sum

- Suppose a set of n values, a_0, a_1, \dots, a_{n-1} , and an associative binary operator \oplus .
 - **Reduction**: computing $a_0 \oplus a_1 \oplus \dots \oplus a_{n-1}$.
 - Addition is an associative binary operator.
 - finding the sum, $a_0 + a_1 + \dots + a_{n-1}$, is a reduction.
- Let's consider a parallel reduction algorithm for summing up n values.
 - Sequential algorithm needs $n-1$ additions.

Parallel Reduction Design (1/5)

■ Partitioning

- Let's divide the list of n values into n pieces
 - A problem is divided as finely as possible $\rightarrow n$ tasks.

■ Communication

- Suppose tasks **A** and **B**.
 - **A** and **B** must be connected via a channel.
 - **A** cannot access a value stored in the memory of **B**.
 - A channel from **A** to **B** is required for **B** to compute the sum of two values.
 - In T/C model, each channel is unidirectional.


Parallel Reduction Design (2/5)

■ Communication (cont'd)

- Let λ and χ be the comm. and comp. times, resp.
 - Case 1: $(n-1)(\lambda + \chi)$
 - Case 2: $(n/2-1)(\lambda + \chi) + (\lambda + \chi) = (n/2)(\lambda + \chi)$
 - Case 3: $(n/4-1)(\lambda + \chi) + 2(\lambda + \chi) = (n/4+1)(\lambda + \chi)$


for addition of two subtotals

Case 1




One task receives $n-1$ values, and performs all the additions.

Case 2



Two tasks work together. The time is cut nearly in half. But, one more comm./comp. step is needed.

Case 3




Four tasks cooperate. These are combined in two more comm./comp. steps.

Parallel Reduction Design (3/5)

■ Communication (cont'd)


- Reduction is performed in **log n** comm. steps.
 - using **binomial trees**.
 - If n is not a power of 2, it is done in $\lfloor \log n \rfloor + 1$ steps.
 - For example, 4 communication steps are needed for reduction of 16 tasks.



Parallel Reduction Design (4/5)

■ Agglomeration and Mapping

- Let's map n tasks to p processors
 - each of n and p is a power of 2 and $p \ll n$.
- The number of tasks is static, computations per task are trivial, and comm. pattern is regular.
 - Agglomerate tasks to minimize communications.
Create one agglomerated task per processor.
(See the tree of mapping strategies)



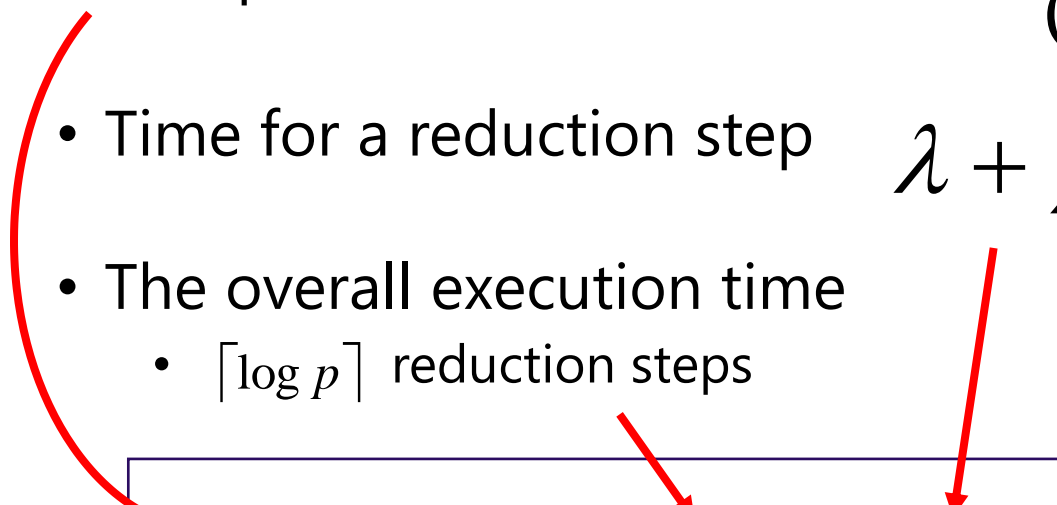
Parallel Reduction Design (5/5)

■ Performance Analysis

- χ : time needed for the binary operation.
- λ : time needed for communication via a channel.
- Computation time for subtotals $(\lceil n / p \rceil - 1)\chi$

- Time for a reduction step $\lambda + \chi$


- The overall execution time
 - $\lceil \log p \rceil$ reduction steps


$$(\lceil n / p \rceil - 1)\chi + \lceil \log p \rceil (\lambda + \chi)$$

The n-Body Problem

■ Newtonian n-body simulation

- Straightforward sequential algorithms have time complexity $O(n^2)$ per iteration, where n is # particles.
 - Of course, there are many better sequential algorithms, though.
- Let's design a parallel algorithm!
 - Simulating the motion of n particles with various masses in a 2-dimensional space.



n -Body Simulation Design (1/3)

■ Partitioning

- Divide the problem as finely as possible $\rightarrow n$ tasks


■ Communication

- **Gather** operation:

- Global communication for **a single task** to collect data items distributed among other tasks.

- **All-gather** operation:


- Global communication for **every task** to collect data items distributed among other tasks.



n-Body Simulation Design (2/3)

■ Communication (cont'd)

- An all-gather operation is needed to update the location and velocity of every particle.
 - Does it need $n-1$ communication steps?
 - No. Remember parallel reduction!
 - → **$\log p$** incoming channels and **$\log p$** outgoing channels for every task



For 4 tasks, $\log 4 = 2$ incoming and outgoing communications are

n-Body Simulation Design (3/3)

■ Agglomeration and Mapping

- Generally, $n \gg p$.
- So each takes one agglomerated task of n/p particles.

■ Performance Analysis

- β : bandwidth (data items sent in one unit of time)
- λ : latency (time to initiate a message)
- χ : time for gravitational force computation
- The communication time per iteration:

$$\sum_{i=1}^{\log p} \left(\lambda + \frac{2^{i-1} n / p}{\beta} \right) = \lambda \log p + \frac{n(p-1)}{\beta p}$$

- The overall execution time per iteration:

$$\lambda \log p + \frac{n(p-1)}{\beta p} + \chi(n/p)$$


Other communication patterns

■ Most programs input and output data.

- Let's add I/O channels to T/C model.
 - No new task = assigning additional duties to a task (I/O task).

■ Scatter operation:

- Global communication like a gather operation in reverse.
- Suppose that n -body simulation, in which I/O task (=task 0) reads n particles from a file.
- Then, I/O task has to send n/p particles to each of the others.(=scatter)



Summary

■ Job Level Parallelism

- What is Job? → A unit of work from user's point of view
- Job Scheduling → substantial performance improvement

■ Parallel Algorithm Design

- Task/Channel Model
 - Tasks and Channels
 - Synchronous and Asynchronous
- Foster's Design Methodology
 - Partition
 - Communication
 - Agglomeration
 - Mapping
- Case Studies
 - Some global communication patterns
 - Gather and Scatter operations