



TOHOKU  
UNIVERSITY



Cyberscience  
Center

# High Performance Computing

## 高性能計算論

Volume 4

Cyberscience Center, Tohoku Univ

**Hiroyuki Takizawa**

**<takizawa@tohoku.ac.jp>**

# What you learnt so far (1/2)

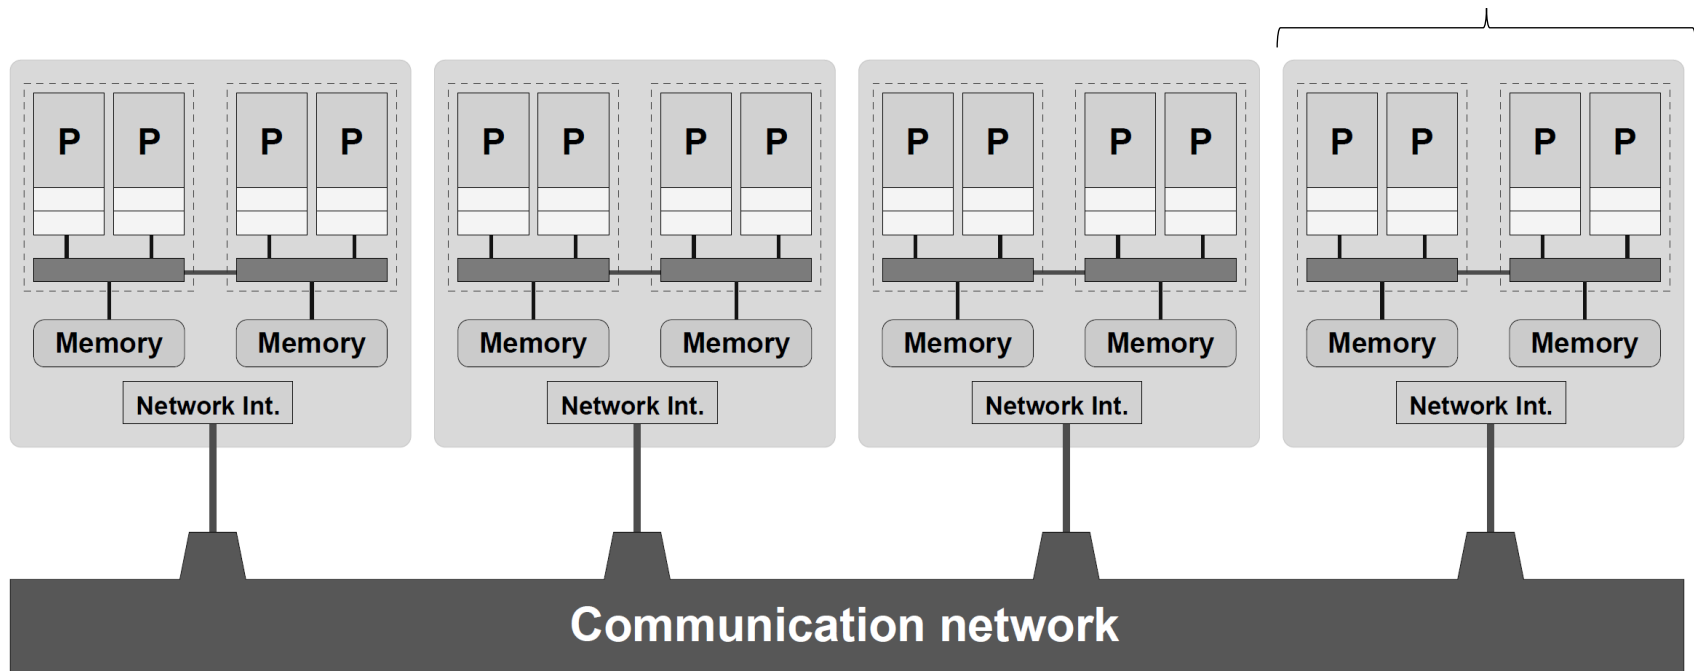
## ■ Parallel Computers

- Vector/SIMD
- Shared-memory computers
- Distributed-memory computers
- Hierarchical (hybrid) systems
- Networks

# Hybrid System

- Large-scale parallel computers  
= mixture of shared and distrib.-parallel.

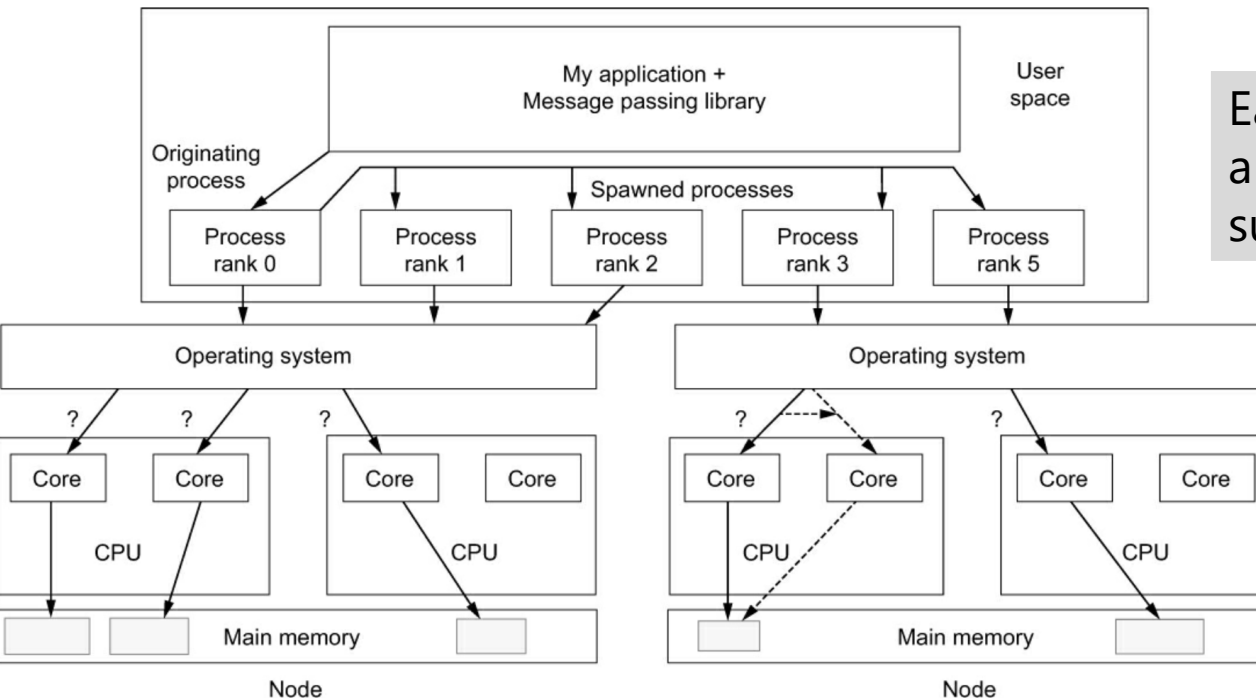
One OS instance manages a node.



In addition, each node may have accelerators such as GPUs.

# Software Overview

When a program is launched, several kinds of resources such as CPU time and memory are allocated for the execution. A unit of allocated resource is called a **process**. An **application** (user program) is executed by multiple processes.



Each process is assigned to a node, which runs the **OS**, such as **Linux**.

The OS on each node decides the core(s) to execute the process. A process can be executed by multiple cores. The execution sequence on each core is called a **thread**.

# What you learnt so far (2/2)

## ■ Job Level Parallelism

- What is Job?
- Job Scheduling

## ■ Parallel Algorithm Design

- Task/Channel Model
- Foster's Design Methodology
- Communication Patterns


# How to run a program on HPC?

## ■ Batch Job

- **A unit of work from user's point of view**
  - Submitted to an HPC System
- A job is usually a batch of tasks
  - Task is a unit of work for a computer


## ■ Front-end Server and Compute Nodes

- Users can log-in to the front-end server, but not directly to the compute nodes
- How to run a job on compute nodes?
  - User submit a job from the front-end server for job execution.



# Job Scheduling

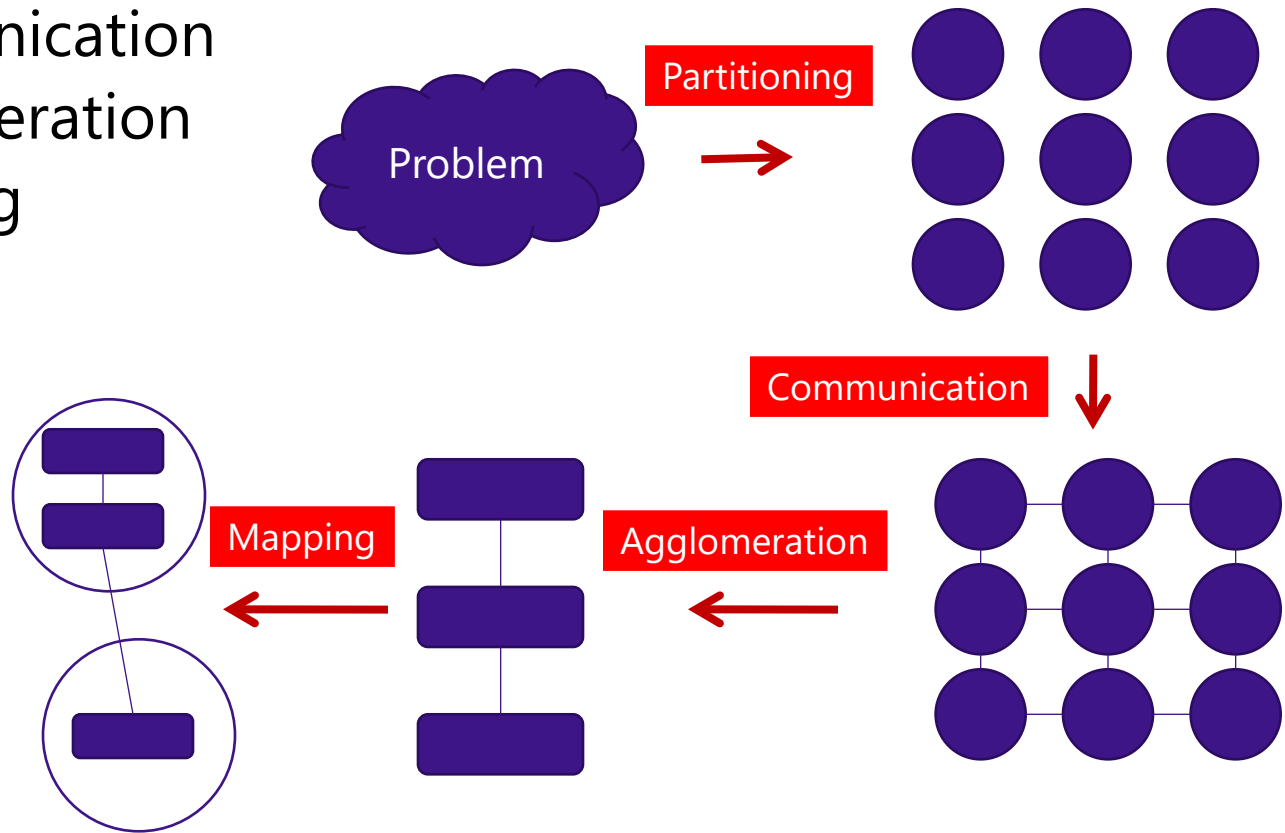
- Decide **where** and **when** a job is executed
  - Necessary for efficient use of shared resources
  - The most basic policy is **First Come First Serve (FCFS)** policy.



# Foster's Design Methodology

## ■ Four steps for designing parallel algorithms

- Partitioning
- Communication
- Agglomeration
- Mapping





# Today's Topic

## ■ Introduction to MPI Programming

- A minimal MPI program
- Performing common communication patterns with collective MPI calls
- Peer-to-peer communication for data exchange
- Benchmarking

You may use any software for programming.  
For example, I use **Visual Studio Code with Remote SSH ext.**

# Programming Parallel Computers

## ■ Four distinct paths (McGraw and Axelrod, 1998)

- Extend an existing compiler
  - Translate a sequential programs into parallel ones
- Extend an existing language
  - Provide new operations to express parallelism
- A new parallel language layer
  - Added on top of an existing sequential language
- A totally new language and compiler system
  - Fortran90, High Performance Fortran, and C\*


## ■ C with MPI and/or OpenMP

- An existing language with low-level constructs

# How to use multiple nodes?

## ■ Parallel Computer = multiple nodes

- Nodes are connected via high-speed network
- Parallel file system is shared by nodes in most cases
- Each node is managed by an OS instance.
  - OS on each node allocates a set of computing resources (= **process**) for program execution.
  - **A process must be created on every node.**



# What we need?

- Parallel computing with **multiple nodes**  
= Parallel computing with **multiple processes**
  - At least one process on every node
    - Otherwise, there is no available resource on the node.

Program launching mechanism

- Data communication among processes
  - Otherwise, each process cannot access data of others.

Data exchange among nodes (= **Message Passing**)

- Synchronization among processes
  - Otherwise, each process can run only independently.

Synchronization and/or blocking communication

# What is MPI?

## ■ Message Passing Interface (MPI)

- Interface for parallel programs with message passing.
  - Multiple processes (**MPI processes**) run on a parallel computer.
    - Each MPI process has its own memory space.
    - MPI processes can pass their data to others if necessary.
  - MPI defines only the interface (not the implementation).
    - MPI defines how each MPI function should work.
    - We do not need to care about the implementation nor internal behaviors.

## ■ Major MPI Implementations

- MPICH (<http://www.mpich.org/>)
- Open MPI (<http://www.open-mpi.org/>)
- MVAPICH (<https://mvapich.cse.ohio-state.edu/>)
  - Computer vendors also provide their own implementations

# Naming Conventions

## ■ C Function Names

- C function name is a format of **MPI\_Xxx\_yyy**
  - MPI\_Init, MPI\_Finalize, MPI\_Send, MPI\_Recv, ...
- The error code is returned
  - **ierr** = MPI\_Comm\_size (MPI\_COMM\_WORLD, &a);

## ■ Keywords (Macro Definitions)

- Keyword is a format of **MPI\_AAA\_BBB**
  - MPI\_COMM\_WORLD, MPI\_INT, MPI\_SUM, ...

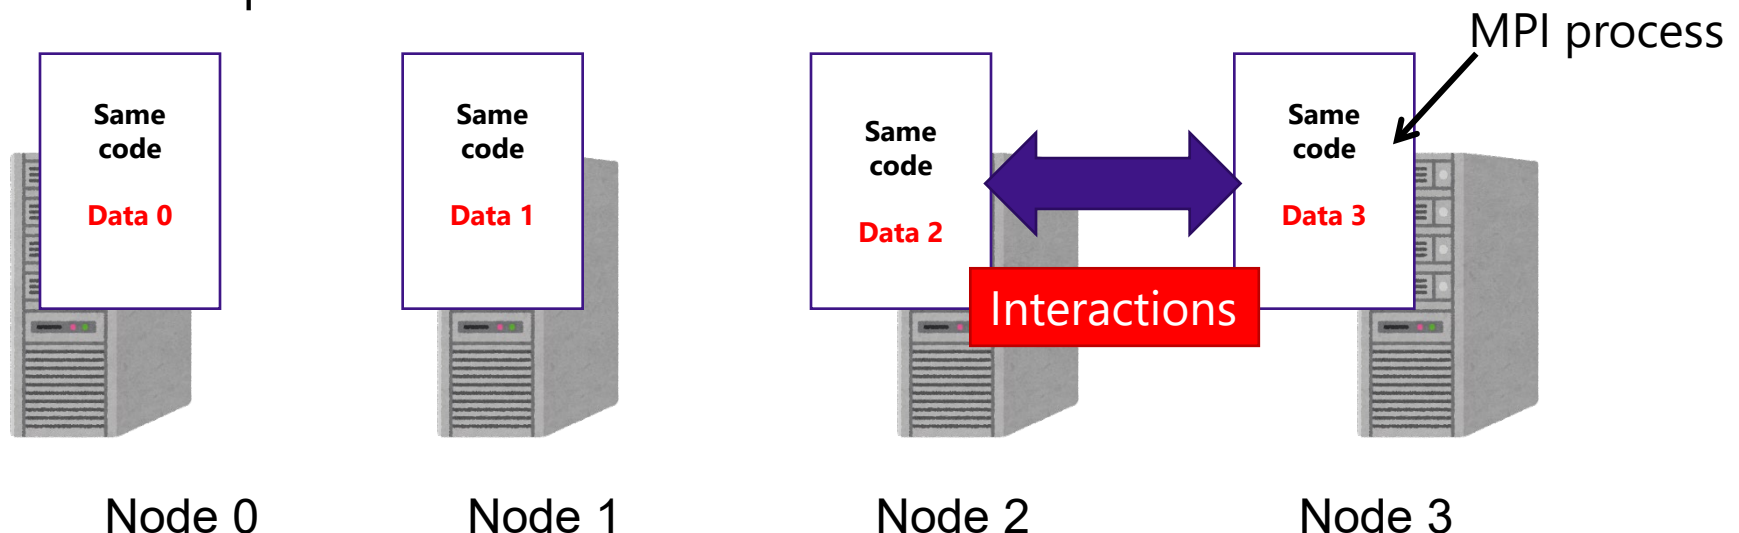
## ■ Fortran Function Names

- Fortran is **case-insensitive**
  - MPI\_INIT, mpi\_init, mpi\_send, MPI\_RECV, ...
- Error code is passed via an argument
  - call mpi\_comm\_size (MPI\_COMM\_WORLD, a, **ierr**)

# MPI Programming Model

## ■ Single-Program Multiple-Data (SPMD)

- The same program is executed on multiple nodes.
  - Programmers write only one source code
  - An **MPI process** is launched on each node.
- Each MPI operates on different data
  - MPI is to explicitly describe interactions among MPI processes



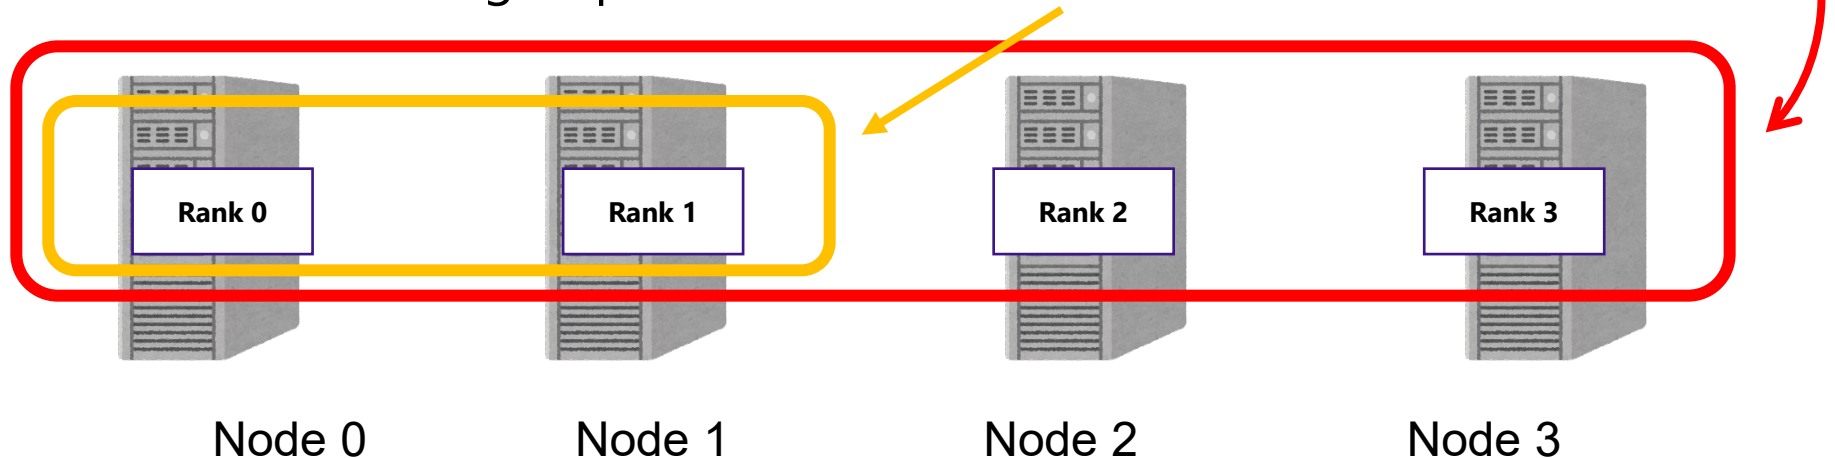
# MPI Programming Model

## ■ MPI Rank = ID number of each process

- Each MPI process has a unique ID number (=rank)

## ■ Communicator = Group of MPI processes

- MPI processes can be grouped into a communicator
- All MPI processes are included in **MPI\_COMM\_WORLD**.
- Each MPI process may join multiple communicators.
  - A subgroup can be defined as **another communicator**.





# Minimal MPI Program

```
#include <mpi.h>
```

```
int main(int argc, char** argv)
{
    MPI_Init(&argc, &argv);
    MPI_Finalize();
    return 0;
}
```

Each MPI function name is a format of **MPI\_Xxxx\_yyyy**

# Compiling and Running MPI program

## ■ Compiling an MPI program


- `mpicc -o min-mpi min-mpi.c`
  - `mpicc` is a compiler command
  - C code program named "min-mpi.c" is converted to an executable "min-mpi".

## ■ Launching an MPI program

- `mpirun -np 1 ./min-mpi`
  - Run "min-mpi" on a single node
- `mpirun -np 4 ./min-mpi`
  - Run "min-mpi" on 4 nodes

# Parallel Reduction

## ■ Finding the sum of $n$ values



# MPI functions

## ■ MPI\_Init & MPI\_Finalize

- Startup and cleanup an MPI runtime environment
  - MPI\_Init must be called before calling any other MPI functions
  - MPI\_Finalize must be called at the end of program execution

## ■ MPI\_Comm\_size

- Retrieves the number of processes in a **communicator**.
  - MPI processes can be grouped into communicators.
  - **MPI\_COMM\_WORLD** is the default communicator that contains all MPI processes working together.

## ■ MPI\_Comm\_rank

- Retrieves the **rank** of the process.
  - Rank is an ID assigned to each process.
  - MPI process can operate different data by using the rank.

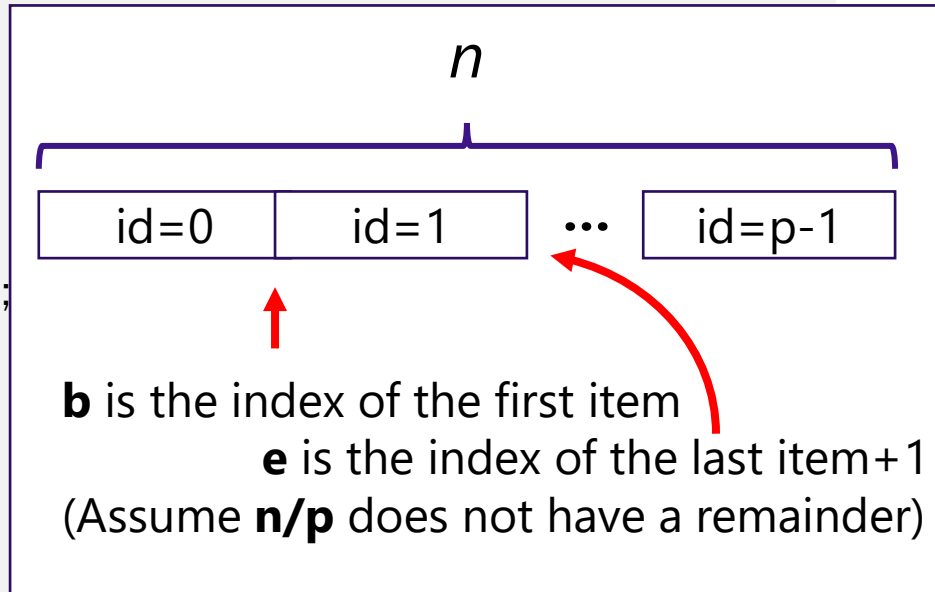
# Sample Code ver. 1

```
#include <mpi.h> /* MPI functions */
#include <stdio.h> /* fprintf */

int main(int argc, char* argv[])
{
    int i, id, p, b, e, s=0, n=10000;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &id);
    MPI_Comm_size(MPI_COMM_WORLD, &p);
    b = (n/p)*id;
    e = (n/p)*(id+1);

    for(i=b; i<e; i++)
        s += i;
    fprintf(stderr, "Process %d is done. ", id);
    MPI_Finalize();
    return 0;
}
```



# Compiling and Running MPI programs

## ■ Compiling MPI programs

- `mpicc -o sum1 sum1.c`
  - `mpicc` is a compiler command
  - C code "sum1.c" is converted to an executable "sum1".

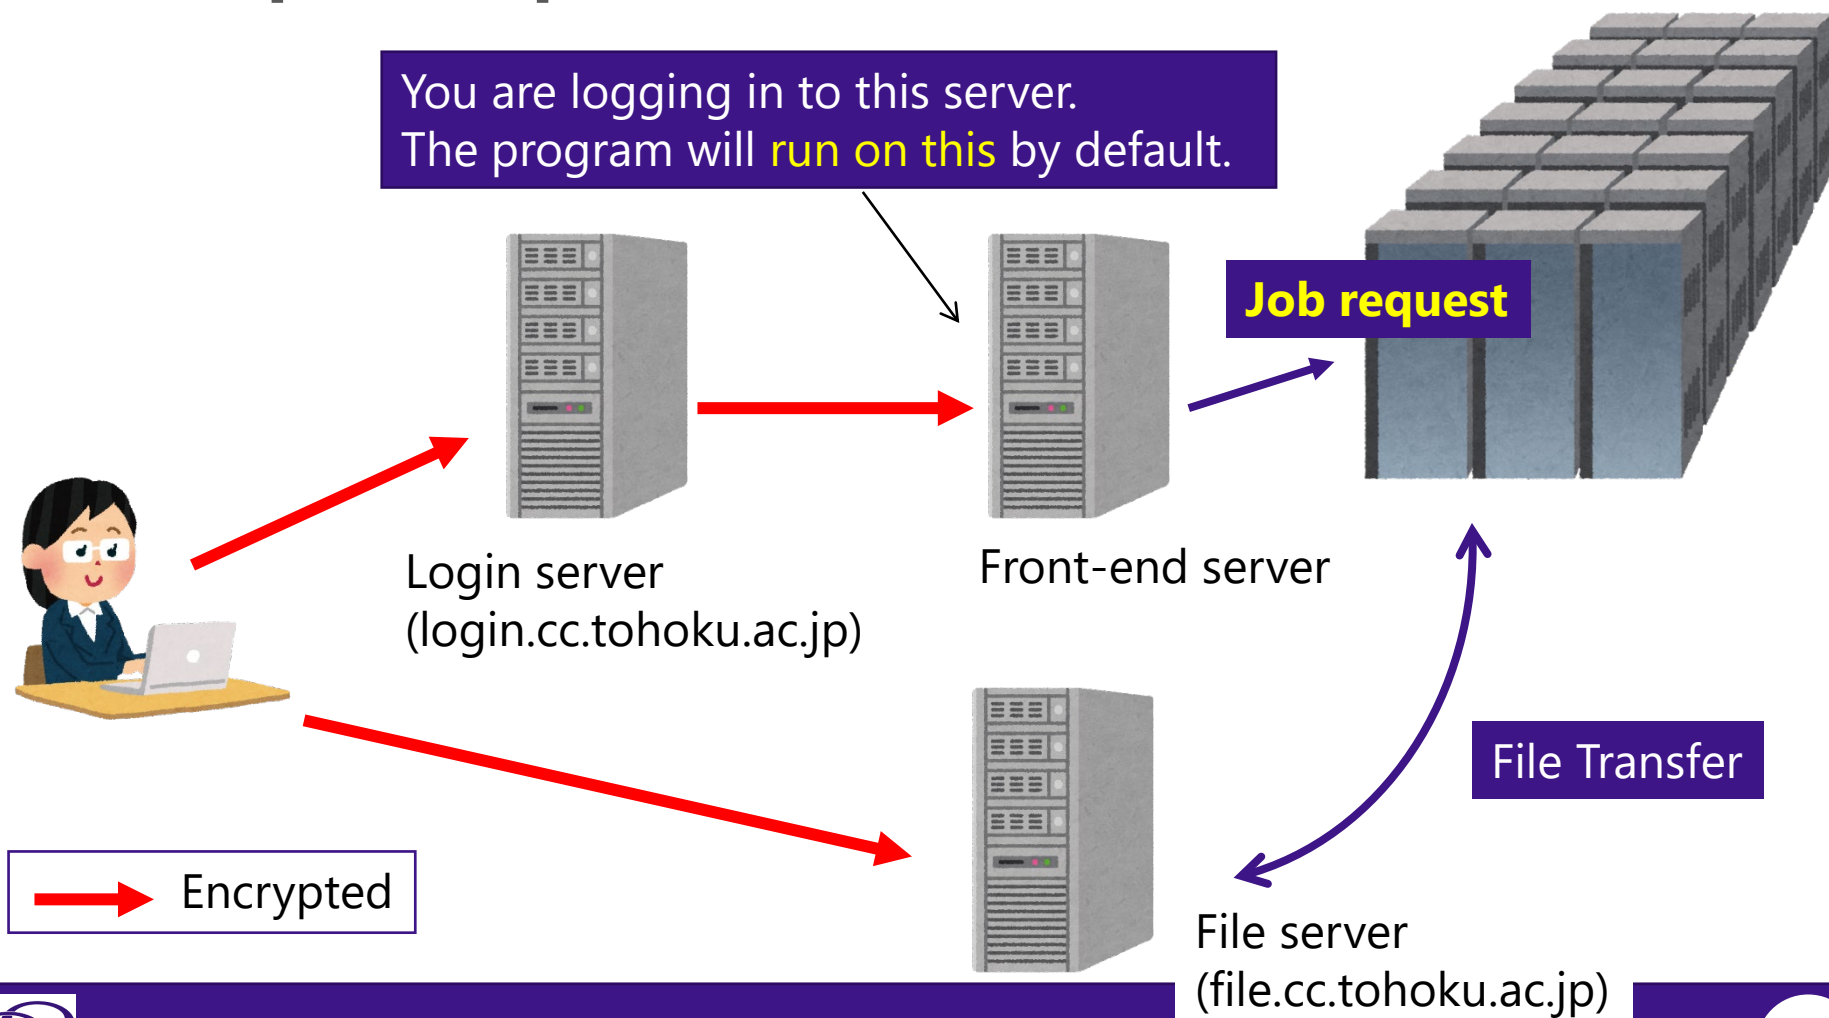
## ■ Running MPI programs

- `mpirun -np 1 ./sum1`
  - Run "sum1" on a single node
- `mpirun -np 4 ./sum1`
  - Run "sum1" on 4 nodes

# Why Job Submission Needed?

## ■ Supercomputer AOBAs = **Shared** Resource

You are logging in to this server.  
The program will **run on this** by default.



# Job Submission


**lx\_edu** is dedicated to this class.  
Use **lx** usually.

## ■ Write a shell script file (run.sh)

```
#!/bin/sh -  
#PBS -q lx_edu  
#PBS -l elapstim_req=0:01:00  
cd $PBS_O_WORKDIR  
mpirun -np 4 ./sum1
```

This job will be executed on **AOBA-B**.

Waiting Queue named "lx\_edu"  
associated with AOBA-B



AOBA-B

AOBA-A

## ■ Submit it to job scheduler

```
qsub run.sh
```

Job submission

Your job will be appended at the end of the queue.

```
qstat
```

Check the status



## ■ Get the result

```
run.sh.eXXXXXX
```

stderr

```
run.sh.oXXXXXX
```

stdout

These files will be created in  
the same directory.  
xxxxx is the job ID. (5 digits)



# What happens?


```
mpirun -np 4 sum1
```

Process 1 is done.  
Process 0 is done.  
Process 3 is done.  
Process 2 is done.

```
less run.sh.exxxxx
```

The content of run.sh.exxxxx

Each process asynchronously works.  
So the messages may differ every  
time they run.



A process of Rank *id* computes a subtotal  
= There is no process that has the total sum.

# MPI Communications

## ■ **Collective** communications

- MPI implementation offers optimized implementations of typical communication patterns
  - Reduction, gather, scatter, broadcast...

## ■ **Peer-to-peer** communications

- In MPI, any pair of two MPI processes can communicate with each other.
  - Send/recv, Isend/irecv, Get/put,

# Sample Code ver. 2

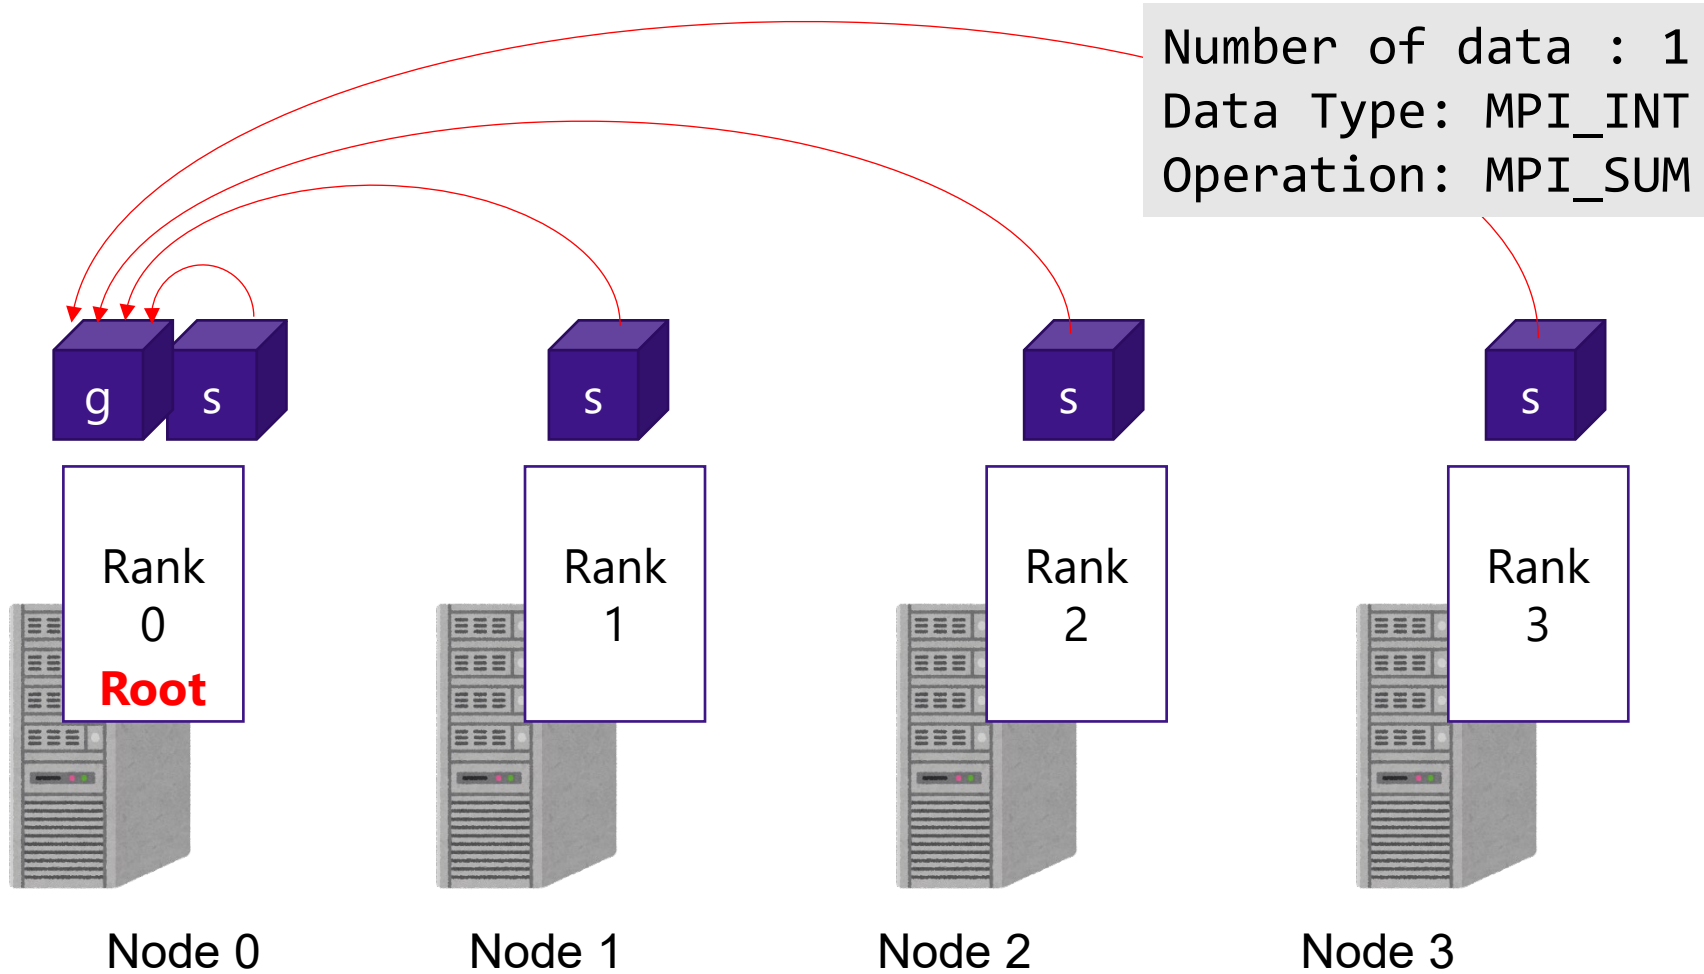
```
#include <mpi.h>    /* MPI functions */
#include <stdio.h>   /* fprintf      */

int main(int argc, char* argv[])
{
    int i, id, p, b, e, s=0, n=10000;
    int g;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &id);
    MPI_Comm_size(MPI_COMM_WORLD, &p);
    b = (n/p)*id;
    e = (n/p)*(id+1);

    for(i=b;i<e;i++)
        s += i;
    MPI_Reduce(&s,&g,1,MPI_INT,MPI_SUM,0,MPI_COMM_WORLD);
    if (id==0)
        fprintf(stderr, "Total=%d\n", g);
    MPI_Finalize();
    return 0;
}
```

# How MPI\_Reduce Works




# Collective Communication

## ■ MPI\_Reduce

- One process becomes a **root** process of a reduction operation, gathering values from all the other processes in a communicator, and summing them up (or other operations).
- Suppose each process has an integer value **s**, and Process 0 calculates the sum of the values, **g**. Then, the reduction operation is

```
MPI_Reduce(&s,&g,1,MPI_INT,MPI_SUM,0,MPI_COMM_WORLD);
```




# Data Types

Name	C Type
MPI_CHAR	signed char
MPI_DOUBLE	double
MPI_FLOAT	float
MPI_INT	int
MPI_LONG	long
MPI_LONG_DOUBLE	long double
MPI_SHORT	short
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long
MPI_UNSIGNED_SHORT	unsigned short

# Reduction Operators

Name	Meaning
MPI_BAND	Bitwise AND
MPI_BOR	Bitwise OR
MPI_BXOR	Bitwise eXclusive OR (XOR)
MPI_LAND	Logical AND
MPI_LOR	Logical OR
MPI_LXOR	Logical eXclusive OR (XOR)
MPI_MAX	Maximum
MPI_MAXLOC	Maximum and its location
MPI_MIN	Minimum
MPI_MINLOC	Minimum and its location
MPI_PROD	Product
MPI_SUM	Sum

# Parallel Reduction






# Parallel Reduction (cont'd)

Optimized communication implementation may be able to reduce the communication cost.

2 steps



# Other Comm. Functions

- **MPI\_Allreduce**: allreduce
- **MPI\_Gather**: gather
- **MPI\_Allgather**: allgather
- **MPI\_Scatter**: scatter
- **MPI\_Bcast**: broadcast
- **MPI\_Send**: send data (blocking)
- **MPI\_Recv**: receive data (blocking)
- **MPI\_Isend**: send data(non-blocking)
- **MPI\_Irecv**: receive data (non-blocking)
- **MPI\_Alltoall**: all-to-all communication


# Reduce and Allreduce

## ■ Reduce operation:

- A single process finally has the result of reduction op.

## ■ All-reduce operation:

- Every process finally has the result of reduction op.





# Gather and Allgather

## ■ Gather operation:

- Global communication for **a single process** to collect data items distributed among others.

## ■ All-gather operation:

- Global communication for **every process** to collect data items distributed among others.





# Scatter and Bcast

## ■ Scatter operation:

- Global communication like a **gather operation in reverse**.
- One MPI process sends **different data** to each of the others.

## ■ Broadcast operation:

- One MPI process sends **the same data** to the others.



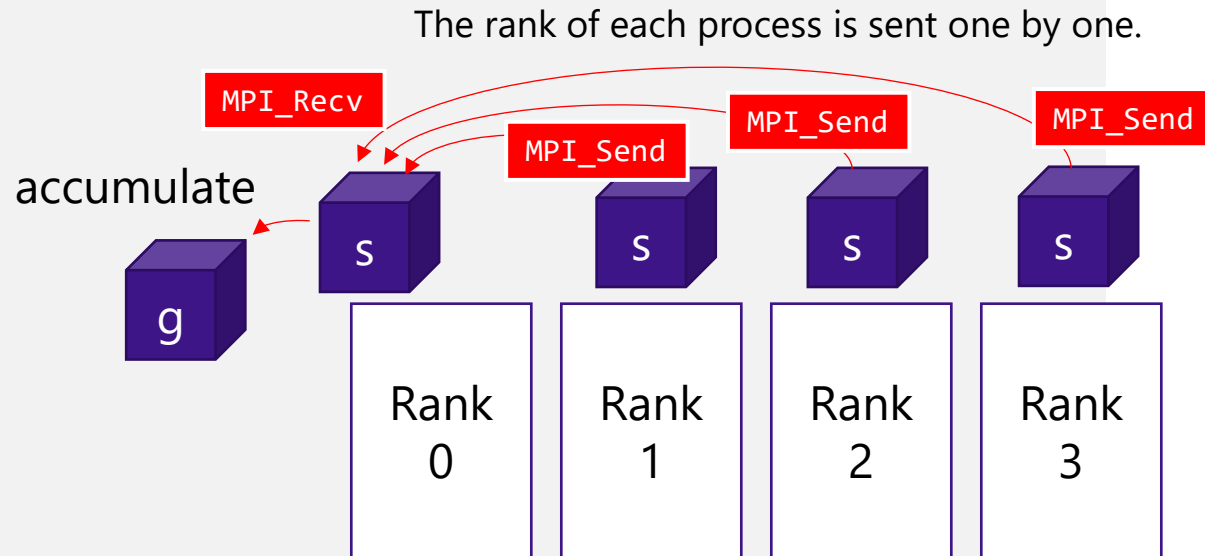
# Sample Code ver. 3

```
#include <mpi.h> /* MPI functions */
#include <stdio.h> /* fprintf */

int main(int argc, char* argv[])
{
    int i, id, p, b, e, s=0, n=10000;
    int g;
    MPI_Status status;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &id);
    MPI_Comm_size(MPI_COMM_WORLD, &p);
    b = (n/p)*id;
    e = (n/p)*(id+1);

    for(i=b;i<e;i++)
        s += i;
    if(id==0) { /* MPI process 0 */
        g = s;
        for(i=1;i<p;i++) { /* receiving values from the others */
            MPI_Recv( &s, 1, MPI_INT, i, 0, MPI_COMM_WORLD, &status);
            g += s;
        }
        printf("The sum is %d.\n", g);
    }
    else { /* The other processes */
        /* sending data to MPI process 0 */
        MPI_Send( &s, 1, MPI_INT, 0, 0, MPI_COMM_WORLD);
    }
    MPI_Finalize();
    return 0;
}
```




# P2P Communication

## ■ MPI\_Send and MPI\_Recv

- The most basic MPI functions for peer-to-peer comm.
  - These functions return after the communication is completed
  - **MPI\_Send**: Sending data to another MPI process
  - **MPI\_Recv**: Receiving data from another MPI process
- Blocking communication
  - These functions return after the communication is completed
  - **MPI\_Isend** and **MPI\_Irecv** are their non-blocking version.

```
MPI_Recv( &j, 1, MPI_INT, i, 0, MPI_COMM_WORLD, &status);
```




When transferring multiple data items at once, a different tag is used to identify each item.

# P2P Communication

## ■ MPI\_Send and MPI\_Recv

- The most basic MPI functions for peer-to-peer comm.
  - **MPI\_Send**: Sending data to another MPI process
  - **MPI\_Recv**: Receiving data from another MPI process
- Blocking communication
  - These functions return after the communication is completed
  - **MPI\_Isend** and **MPI\_Irecv** are their non-blocking version.

```
MPI_Send( &j, 1, MPI_INT, 0, 0, MPI_COMM_WORLD);
```





# Deadlocks

```
#include <mpi.h> /* MPI functions */  
#include <stdio.h> /* printf */
```

```
int main(int argc, char* argv[])  
{  
    int i, j, id, p, src, dst;  
    MPI_Status status;
```

```
    MPI_Init(&argc, &argv);  
    MPI_Comm_rank(MPI_COMM_WORLD, &id);  
    MPI_Comm_size(MPI_COMM_WORLD, &p);
```

```
    dst = (id+1)%p;  
    src = id-1<0?p-1:id-1;  
    MPI_Send(&id, 1, MPI_INT, dst, 0, MPI_COMM_WORLD);  
    /* never reach this point */  
    MPI_Recv(&id, 1, MPI_INT, src, 0, MPI_COMM_WORLD, &status);
```

```
    MPI_Finalize();  
    return 0;  
}
```



Every process is sending data, and waiting (blocked) until the receiver receives it.  
= every process is waiting here.  
= no process can receive the data.

**Deadlock**



# Synchronous / Asynchronous

## ■ Synchronous communication


`MPI_Send` & `MPI_Recv`

- The task will wait until the data are sent/received.  
(= the task is **blocked**.)

## ■ Asynchronous communications

`MPI_Isend` & `MPI_Irecv`

- The task is never blocked by sending/receiving data.  
(= an **asynchronous** or **non-blocking** operation)
- A task can initiate a data communication operation and then do another thing until the operation is completed.
- `MPI_Waitall` is a function to explicitly block the task until the data communication is finished.



**Overlapping computation  
and communication!**

# MPI\_Isend & MPI\_Irecv

```
MPI_Request req[2];  
MPI_Status stat[2];
```

```
// Sending a data item is initiated but the task is not blocked  
MPI_Isend(&send_data[0], send_data_count, MPI_DOUBLE,  
          dest, 0, MPI_COMM_WORLD, &req[0]);
```

```
// Receiving a data item is initiated but the task is not blocked  
MPI_Irecv(&recv_data[0], recv_data_count, MPI_DOUBLE,  
          src, 0, MPI_COMM_WORLD, &req[1]);
```

```
// The task is blocked here until the operations associated with  
// req[0] and req[1] are completed.  
// Error codes of req[0] and req[1] are in stat[0] and stat[1]  
MPI_Waitall(2, req, stat);
```

# Benchmarking Performance

## ■ How to measure the elapsed time for parallel processing?

- **MPI\_Wtime**

- If **MPI\_Wtime** is called twice, the different between their return values indicates the elapsed time [sec] between the calls.

```
double etime;  
etime = -MPI_Wtime();  
/* some parallel task */  
etime += MPI_Wtime();  
printf("elapsed time: %lf [sec]¥n", etime);
```

↑  
The time measurement could be inaccurate!!

# Barrier Synchronization

## ■ How to ensure every process calls MPI\_Wtime at the same time?

- **Barrier synchronization**

- No process can proceed beyond it until all processes have reached it.

- **MPI\_Barrier**


- Barrier synchronization of one communicator group.
- Time-consuming, so prevent unnecessary synchronizations.

```
double etime;  
MPI_Barrier(MPI_COMM_WORLD);  
etime = -MPI_Wtime();  
/* some parallel task */  
MPI_Barrier(MPI_COMM_WORLD);  
etime += MPI_Wtime();  
printf("elapsed time: %lf [sec]\n", etime);
```

# Exercise :

## Monte Carlo $\pi$ Calculation

- The area of a circle is  $\pi=3.1415\dots$  if the radius is 1.
- Suppose that a lot of points are randomly generated within the region of  $0 < x < 1$  and  $0 < y < 1$ .
  - The probability that a point is in a circle is  $\pi/4$ .



Numerical integration by Monte Carlo approach

By increasing the number of random points  
(The number of points in the circle)/(the total number of points)  $\rightarrow \pi/4$ .

# Sequential Code

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char* argv[])
{
    int N = 1000;
    int i, total = 0;
    double x, y;

    srand(time(NULL)); /* initialization */
    for(i=0; i<N; i++) {
        /* two random numbers 0<x,y<1 */
        x = (double)rand()/RAND_MAX;
        y = (double)rand()/RAND_MAX;
        if( x*x + y*y < 1 ) {
            total = total+1;
        }
    }
    printf( "pi=%lf\n", 4.0*total/N);
}
```

Statistically speaking, the accuracy will improve by increasing N.

Hint:  
Divide the loop into some pieces, and assign each piece to one MPI process.  
**Use reduction to calculate the total from subtotals.**

Parallelize this code and measure the performance!

# Summary

## ■ Introduction to MPI Programming

- A minimal MPI program
- Performing common communication patterns with collective MPI calls
- Peer-to-peer communication for data exchange
- Blocking and non-blocking operations
- Benchmarking

## ■ T/C model naturally fits MPI paradigm.

- Let's try more complicated parallel program design!