

High Performance Computing (Nov 12)

Question 1

Run the OpenMP program on Slide 16 with changing the number of threads. How does the execution time change with the thread count?

Answer

The code calls the `sleep()` function 16 times. At each function call, the process sleeps for 1 second. As a result, the process is expected to take 16 seconds in the case of single-thread execution. If the process is executed with N threads, each thread simultaneously calls the function and thus sleeps for 1 second at the function call. In theory, hence, the execution time will be $16/N$ seconds.

Sample job script (run.sh)

```
-----  
#!/bin/sh  
#PBS -q lx_edu  
  
cd $PBS_O_WORKDIR  
  
export OMP_NUM_THREADS=1  
time ./a.out  
export OMP_NUM_THREADS=2  
time ./a.out  
export OMP_NUM_THREADS=4  
time ./a.out  
export OMP_NUM_THREADS=8  
time ./a.out  
export OMP_NUM_THREADS=16  
time ./a.out  
-----
```

Execution results (run.sh.eXXXXXX)

```
-----  
real    0m16.004s  
user    0m0.001s
```

```
sys      0m0.000s
```

```
real     0m8.003s
```

```
user     0m0.002s
```

```
sys      0m0.000s
```

```
real     0m4.003s
```

```
user     0m0.002s
```

```
sys      0m0.001s
```

```
real     0m2.003s
```

```
user     0m0.000s
```

```
sys      0m0.007s
```

```
real     0m1.004s
```

```
user     0m0.020s
```

```
sys      0m0.003s
```

The `time` command shows the total execution time (wall clock time) as “real,” the CPU time (spent by the program itself) as “user,” and the execution time spent by the OS as “sys.” The sample program calls the `sleep()` function 16 times to increase the total execution time. But the function call just sleeps the process, and thus does not spend the CPU time or the OS time for the execution.

Each thread sleeps at the same time. As a result, the execution time decreases with the thread count. In this example, the runtime overhead is negligibly small, and the execution time decreases almost perfectly with the thread count.

Question 2

Write a code to check the behavior of reduction code in Slide 22, and paste it.

Answer

A simple example to check the behavior is as follows.

```
-----  
#include <stdio.h>  
#include <unistd.h>
```

```
#define N 256
```

```
int main(int ac, char* av)
{
    int i, j, x;

    x=0;
    #pragma omp parallel for reduction(+:x), private(j)
    for(i=0; i<N; i++){
        for(j=0; j<N; j++){
            x+=j;
        }
    }
    printf("%d\n", x);
    return 0;
}
```

You can check the behavior by enabling and disabling OpenMP directives. It is enabled only when the command line option of `-fopenmp` is given at compilation.

```
$ gcc -fopenmp omp2.c
```

```
$ ./a.out
```

```
8355840
```

```
$ gcc omp2.c
```

```
$ ./a.out
```

```
8355840
```

As you can see, the computation result does not change, meaning that the reduction clause works correctly.