



TOHOKU
UNIVERSITY



Cyberscience
Center

High Performance Computing

高性能計算論

Volume 5

Cyberscience Center, Tohoku Univ

Hiroyuki Takizawa

<takizawa@tohoku.ac.jp>

Class Schedule

1-Oct Introduction to HPC (1)

1-Oct Introduction to HPC (2)

8-Oct Parallel Architectures

8-Oct How to use Supercomputer AOBA

15-Oct Parallel Algorithm Design (1)

15-Oct Parallel Algorithm Design (2)

5-Nov MPI Programming (1)

5-Nov MPI Programming (2)

12-Nov OpenMP Programming (1)

12-Nov OpenMP Programming (2)

26-Nov Hybrid Programming

26-Nov Performance Modeling and Analysis

← Pre-recorded video

Today's Topic

■ Introduction to OpenMP Programming

- How to parallelize a for loop.
- Shared data and private data.
- Task-parallel processing
- Offloading

■ Some keys to achieve better scalability

What's Parallel Computer (1/3)

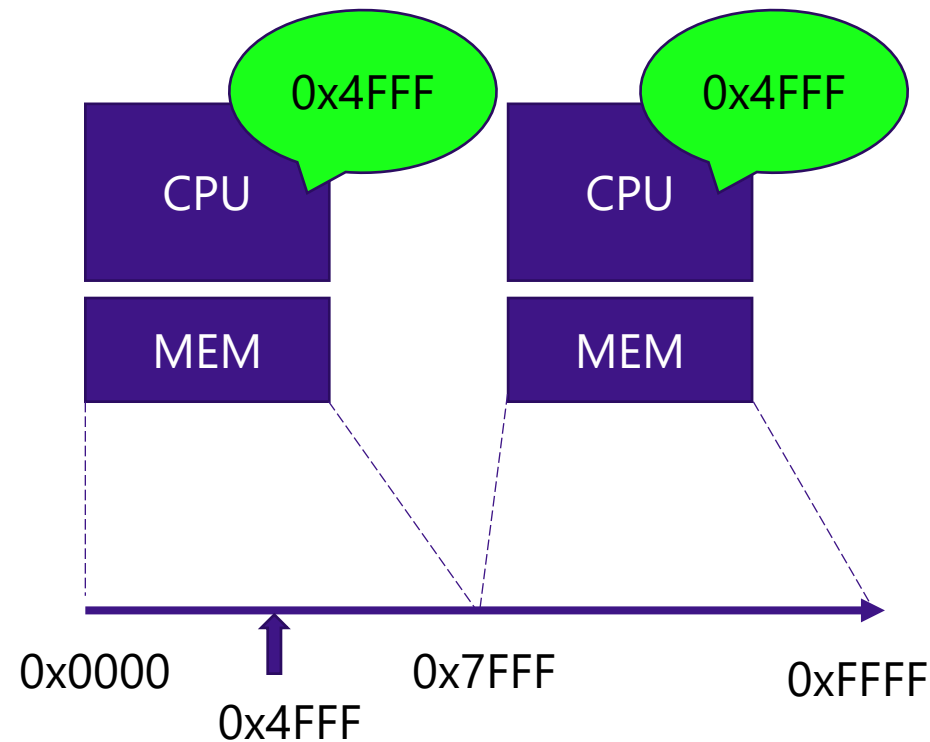
■ Parallel Computer

- A multi-processor computer system supporting parallel programming
- Two major categories of parallel computers
 - Distributed-memory parallel computers
 - Multiple computers and their interconnection network.
 - Employed to build a large-scale system
 - Shared-memory parallel computers
 - Symmetric multi-processor(SMP) and multicore/manycore.
 - Employed by most of current processors.

What's Parallel Computer (2/3)

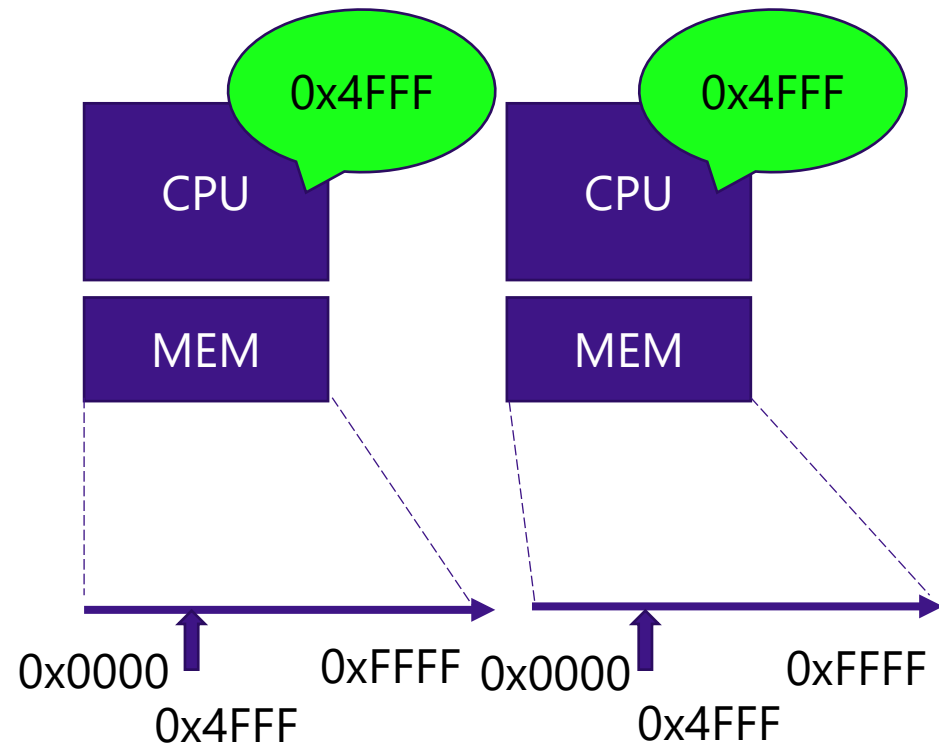
Shared-memory (NUMA)

Each memory device is mapped to a part of the memory space.



Distributed-memory

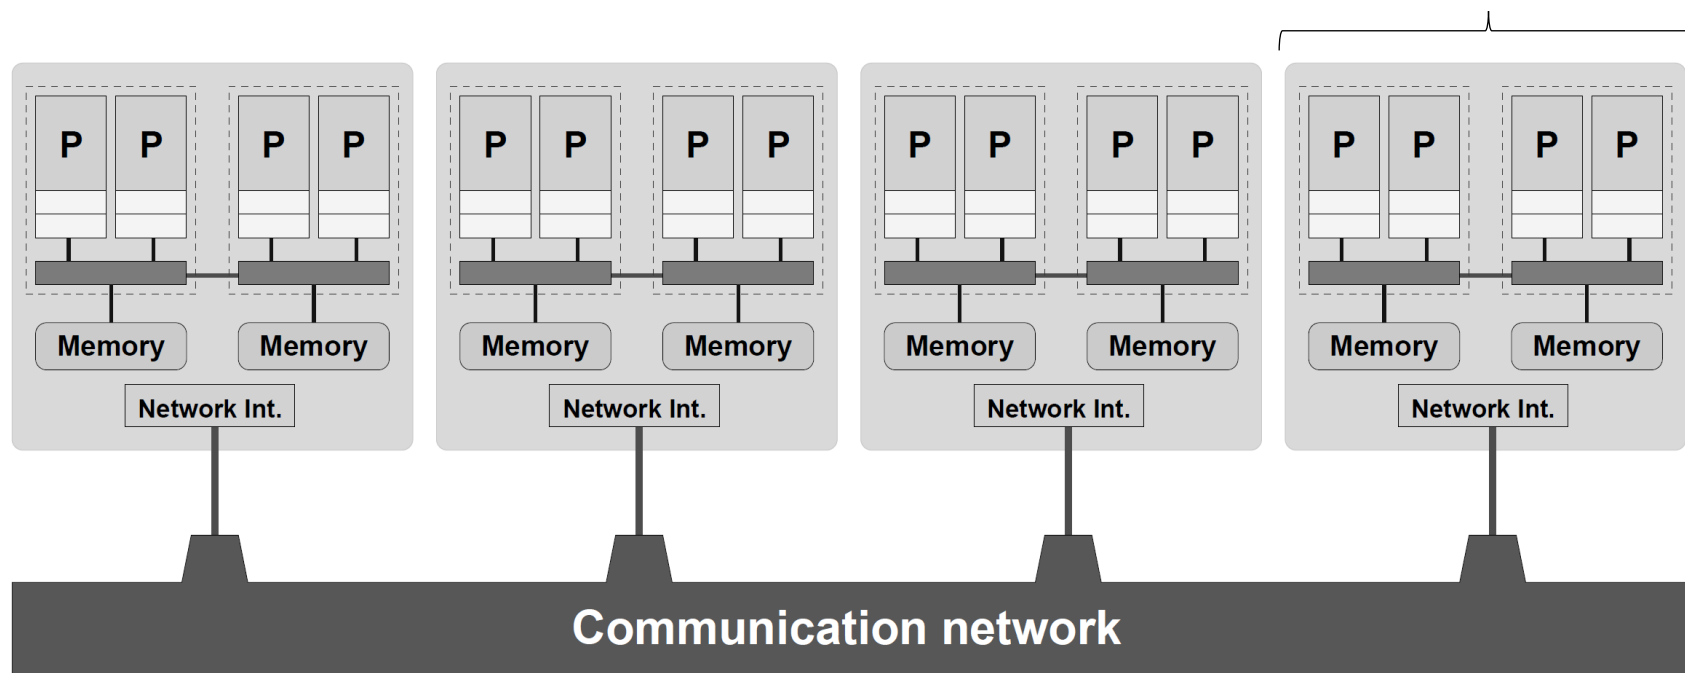
Each memory device has its own memory space.



What's Parallel Computer (3/3)

- Large-scale parallel computers
= mixture of shared and distrib.-parallel.

One OS instance manages a node.



In addition, each node may have accelerators such as GPUs.

What's Parallel Programming

■ Parallel Programming

- Programming in a language that allows you to explicitly indicate parallel portions of a program.
 - Those portions can run concurrently on different processors
- Why required?
 - Automatic parallelization is difficult, even though there have been many studies on parallelizing compilers.
- **OpenMP** and **MPI**
 - Standards for shared- and distributed-memory systems, respectively

Block in C/C++ Language

- In C/C++ language, a **block** is the code between **{** and **}**.

EX) body of a for loop.

```
for (i=0;i<10;i++)  
    printf("hello1¥n");  
    printf("hello2¥n");
```

```
for (i=0;i<10;i++){  
    printf("hello1¥n");  
    printf("hello2¥n");  
}
```

OpenMP is used to specify how to execute a block.

Processes and Threads

■ Both processes and threads are execution flows of a program.

- When a program is launched, OS reserves a set of some computing resources for the execution, a so-called **process**.
 - CPU time
 - Memory space
 - File descriptors
- A **thread** is created inside a process
 - CPU time is assigned to each thread
 - The other resources are shared with the other threads of the process.

Compiler Directives

```
#include <stdio.h>
```

```
#define N 1000
```

```
int main(){
```

```
#ifdef DEBUG
```

```
    printf("debug mode");
```

```
#else
```

```
    printf("normal mode");
```

```
#endif
```

```
    return 0;
```

```
}
```

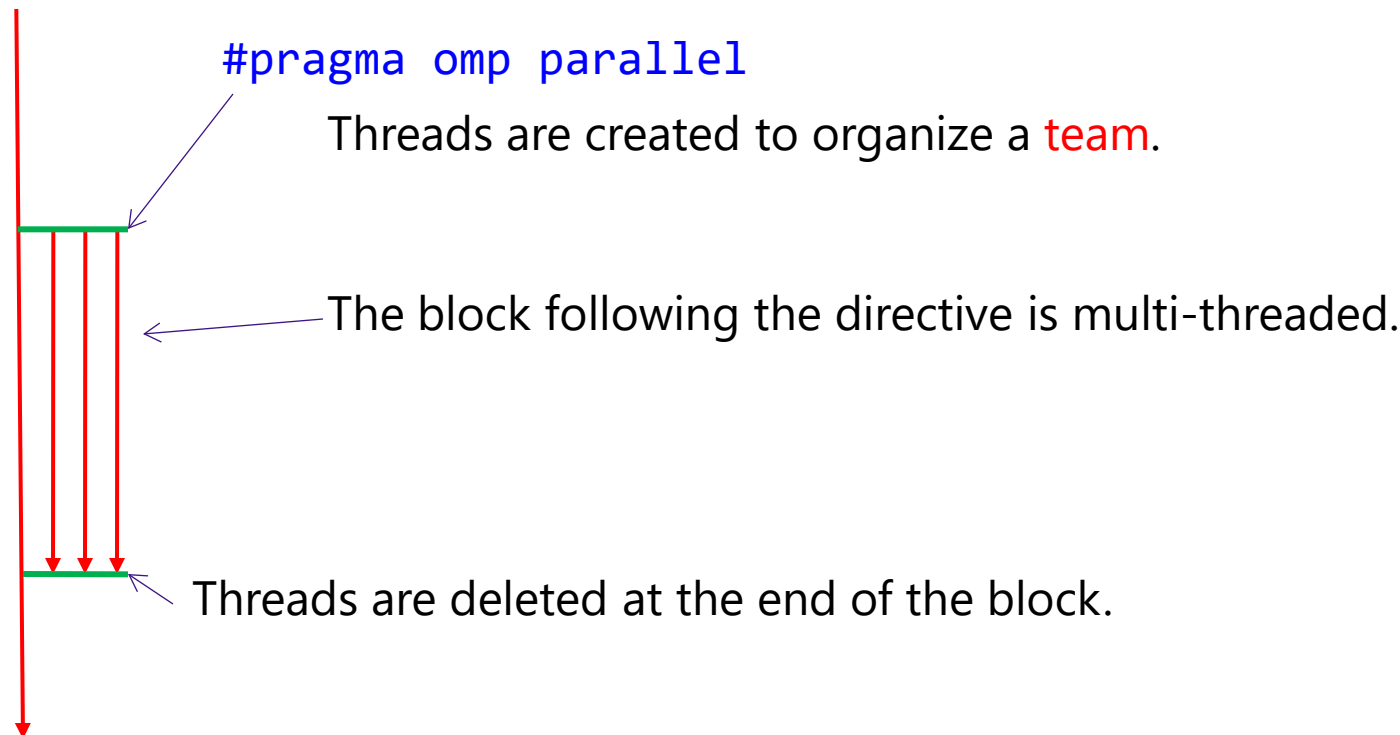
The line starting with # is not a statement, but a compiler directive to specify compiler's behavior.

What's OpenMP?

■ Threads are created/deleted on demand.

- Thread: an execution flow

Fork-join model



Data-Parallel Processing

■ **Work-sharing** by inserting directives into a sequential code.

- C code

```
for(i=0;i<10;i++)  
    a[i] = b[i]*f + c[i];
```

- OpenMP code

```
#pragma omp parallel  
{  
    #pragma omp for  
    for(i=0;i<10;i++)  
        a[i] = b[i]*f + c[i];  
}
```

CPU0

```
a[0] = b[0]*f+c[0]  
a[1] = b[1]*f+c[1]  
a[2] = b[2]*f+c[2]  
a[3] = b[3]*f+c[3]  
a[4] = b[4]*f+c[4]  
a[5] = b[5]*f+c[5]  
a[6] = b[6]*f+c[6]  
a[7] = b[7]*f+c[7]  
a[8] = b[8]*f+c[8]  
a[9] = b[9]*f+c[9]
```

CPU0

```
a[0] = b[0]*f+c[0]  
a[1] = b[1]*f+c[1]  
a[2] = b[2]*f+c[2]  
a[3] = b[3]*f+c[3]  
a[4] = b[4]*f+c[4]
```

CPU1

```
a[5] = b[5]*f+c[5]  
a[6] = b[6]*f+c[6]  
a[7] = b[7]*f+c[7]  
a[8] = b[8]*f+c[8]  
a[9] = b[9]*f+c[9]
```

Data-Parallel Processing

■ **Work-sharing** by inserting directives into a sequential code.

- C code

```
for(i=0;i<10;i++)  
    a[i] = b[i]*f + c[i];
```

- OpenMP code (simplified)

```
#pragma omp parallel for  
for(i=0;i<10;i++)  
    a[i] = b[i]*f + c[i];
```

parallel and **for** can be written together in one line

CPU0

```
a[0] = b[0]*f+c[0]  
a[1] = b[1]*f+c[1]  
a[2] = b[2]*f+c[2]  
a[3] = b[3]*f+c[3]  
a[4] = b[4]*f+c[4]  
a[5] = b[5]*f+c[5]  
a[6] = b[6]*f+c[6]  
a[7] = b[7]*f+c[7]  
a[8] = b[8]*f+c[8]  
a[9] = b[9]*f+c[9]
```

CPU0

```
a[0] = b[0]*f+c[0]  
a[1] = b[1]*f+c[1]  
a[2] = b[2]*f+c[2]  
a[3] = b[3]*f+c[3]  
a[4] = b[4]*f+c[4]
```

CPU1

```
a[5] = b[5]*f+c[5]  
a[6] = b[6]*f+c[6]  
a[7] = b[7]*f+c[7]  
a[8] = b[8]*f+c[8]  
a[9] = b[9]*f+c[9]
```

How many threads?

- Set environment variable, **OMP_NUM_THREADS**, to configure the number of threads for running the program
 - B shell
 - `export OMP_NUM_THREADS=4` ← On AOBA, use this one.
 - C shell
 - `setenv OMP_NUM_THREADS 4`

Compilation and Execution

■ Environment Variable **OMP_NUM_THREADS**

- specifies the number of threads launched in OpenMP
- EX) **export OMP_NUM_THREADS=4**
 - 4 threads are created for a team.

■ **GNU C/C++ Compilers (ver. 4.2 or later)**

- gcc/g++ options source

g++ -fopenmp sample.cc
[command] [options] [source code file]

Exercise: simple code

```
#include <stdio.h>
#include <unistd.h>
int main(int ac, char* av)
{
    int i;
    for(i=0; i<16; i++){
        sleep(1);
    }
    return 0;
}
```

[How to run]

```
% gcc hoge.c
% time ./a.out
```

```
#include <stdio.h>
#include <unistd.h>
int main(int ac, char* av)
{
    int i;
    #pragma omp parallel for
    for(i=0; i<16; i++){
        sleep(1);
    }
    return 0;
}
```

[How to run]

```
% gcc -fopenmp hoge.c
% time ./a.out
```


Shared Variables

■ Threads share a single memory space.

- **All data are shared by threads by default.**

The loop index variable is an exception (not shared).

- EX) Finding the maximum in an array

```
max = 0;
#pragma omp parallel for
for(i=0;i<10;i++){
    if( max < a[i] ) max=a[i];
}
```

NOTE!

Don't write a code like this.

Multiple threads may compete to write values to variable "max."

Race condition



Private Variables and Critical Section

■ Threads share a single memory space.

- How can each thread hold a unique value?

= private variables

clause: optional component to pragma

```
#pragma omp parallel private (tmp)
{
    tmp = 0;
    #pragma omp for
    for(i=0;i<10;i++)
        if(tmp < a[i]) tmp = a[i];

    #pragma omp critical
    {
        if(max<tmp) max = tmp;
    }
}
```


Multi-threaded

Each thread holds a unique value in tmp.
(tmp keeps the maximum value in a subarray)

Critical Section

Once a thread enters the section,
the others cannot enter it.
(See Next Slide)

Private Variables and Critical Section



Even if one thread updates its own tmp, tmp in view from the others does not change.

If one thread updates max, max in view from the others also changes.


While one thread is updating max, we have to inhibit the others from accessing it.

⇒ `#pragma omp critical`

Sections

- If there are multiple sections that can be done in parallel...

```
#pragma omp sections
{
  #pragma omp section
  {
    A
  }
  #pragma omp section
  {
    B
  }
}
```



They are executed in parallel by different threads

Major OpenMP Directives

■ **#pragma omp parallel**

- directs parallelizing the following block

■ **#pragma omp for**

- directs work-sharing the following for loop

■ **#pragma omp sections**

- directs the following block is a set of parallel sections

■ **#pragma omp section**

- directs the following block is a section

■ **#pragma omp critical**

- directs the following block is a critical section
 - Just one thread can enter it

Clauses

■ **shared (default)**

- The variables are shared by threads

■ **private/firstprivate/lastprivate/copyprivate**

- The variables can have different values for different threads. The value can be copied from/to the outside of the block.

■ **nowait**

- Each thread does not wait for the others at the end of the block.

■ **reduction**

- Reduction operations are applied to the variables.

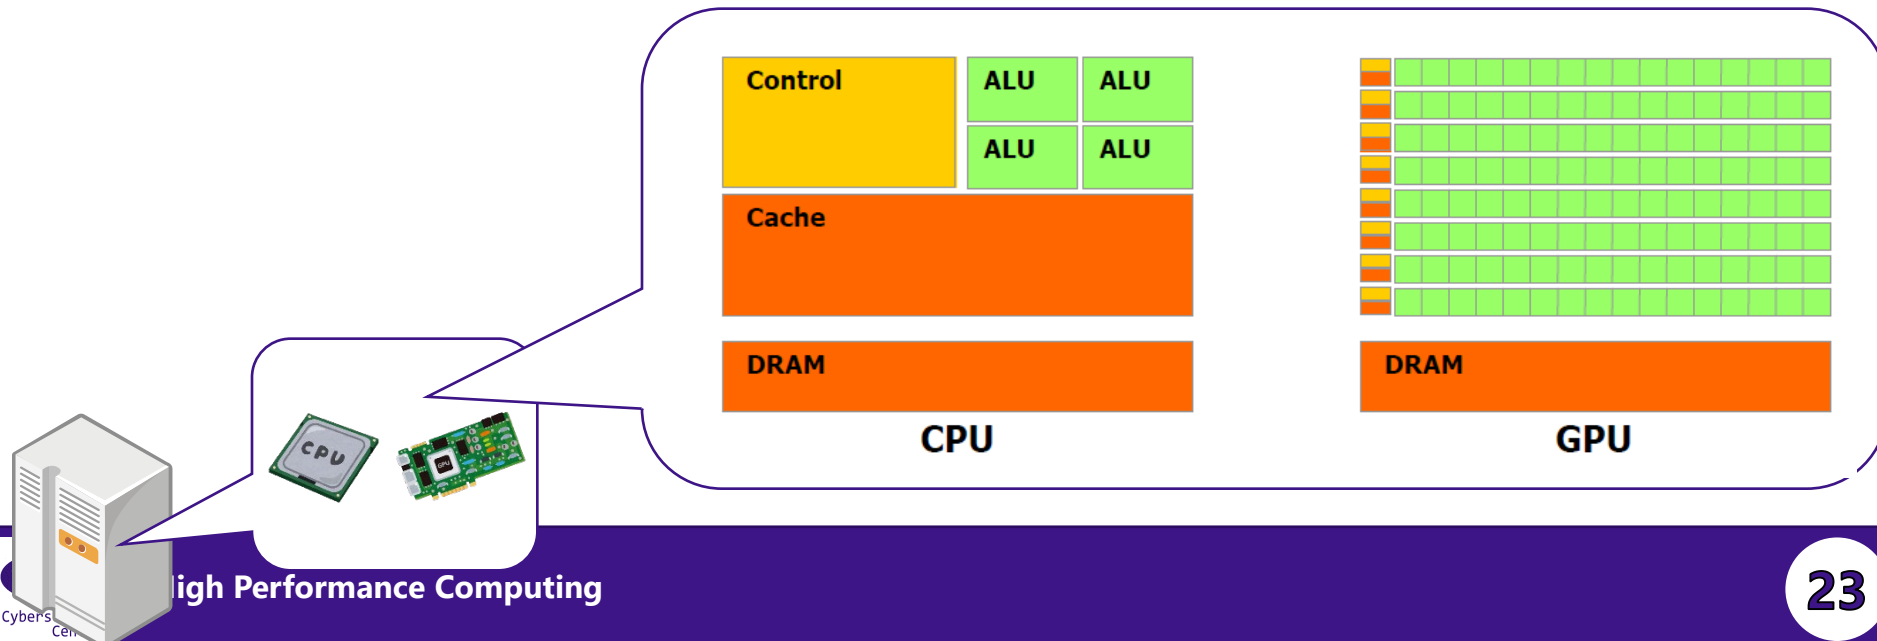
```
#pragma omp for reduction(+:x), private(j)
for(i=0;i<N;i++) {
    for(j=0;j<N;j++) {
        x+=j;
    }
}
```

Homework: Confirm this code fragment works correctly.
(You need to add the prolog and epilog codes to make this work.)

Heterogenous Computing

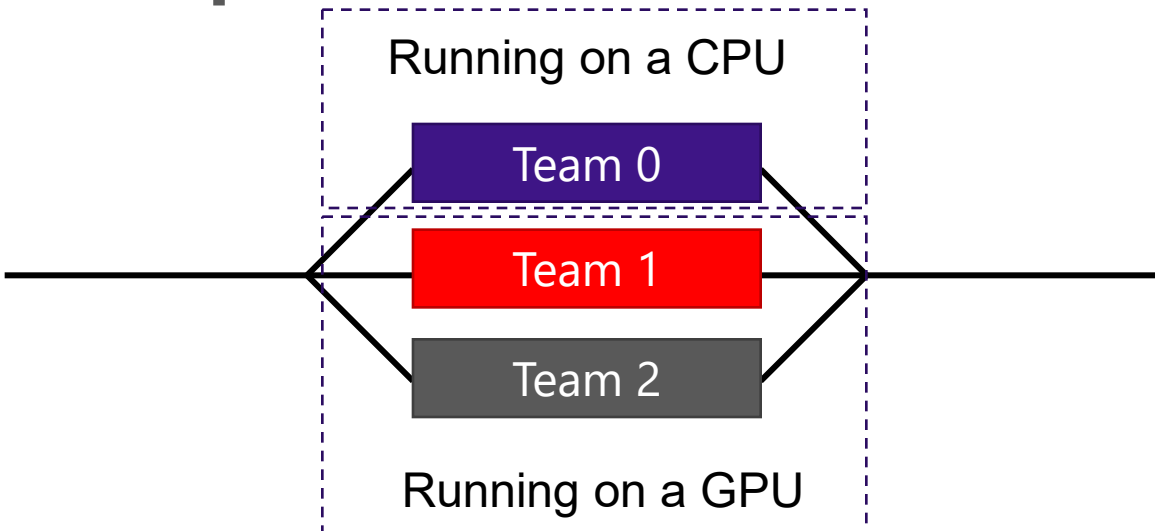
■ CPU and GPU are very different processors.

- Latency-oriented design (=speculative)
 - CPU has a large cache memory and control unit.
- Throughput-oriented design (=parallel)
 - GPUs devote more hardware resources to ALUs.



GPU Programming with OpenMP

- A task can be executed on a different processor



```
#pragma omp target map (alloc:Anew) map(tofrom:A)
{
#pragma omp teams distribute
  for(j=0;j<M;j++){
    #pragma omp parallel for
    for(i=;i<N;i++){
      /* time-consuming data-parallel computation */
    }
  }
}
```


Target Directives

■ #pragma omp target

■ #pragma omp target data Defining only data mapping

- New features of OpenMP 4.0 or later
- The block specified by the target directive can be **offloaded to another processor such as GPU.**
- A map clause is used to send/retrieve data to/from the GPU.

```
/* Array Anew is created, and A is transferred from/to GPU */  
#pragma omp target map(alloc:Anew) map(tofrom:A)  
{  
#pragma omp teams distribute parallel for  
  for(i=;i<N;i++){  
    /* time-consuming data-parallel computation */  
  }  
}
```

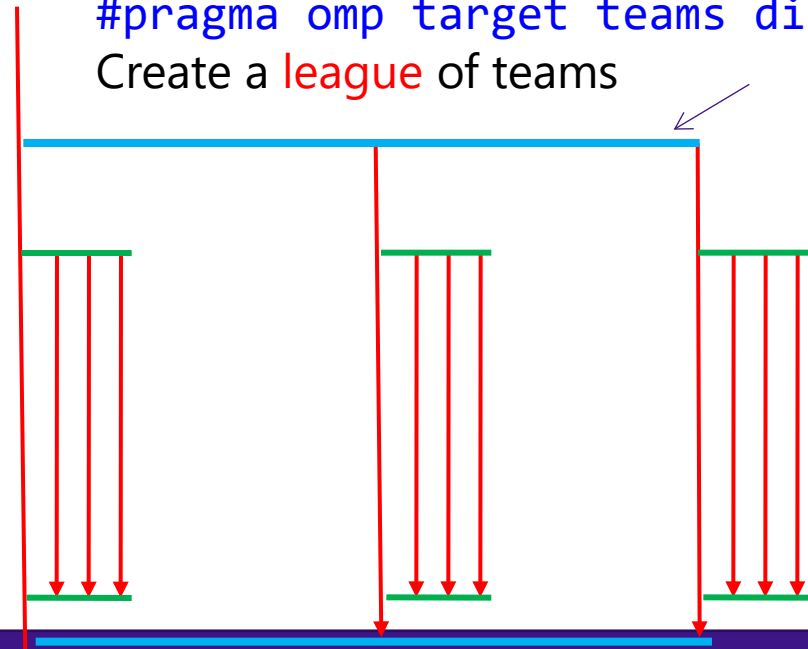
This part is executed on GPU

Target Directives (Cont'd)

- `#pragma omp teams distribute`
 - By definition, all threads in a team have to be synchronizable.
 - A many-core processor needs to run many threads, but the synchronization overhead increases with the number of participating threads.

`#pragma omp target teams distribute`
Create a **league** of teams

`#pragma omp parallel`
Creates a team of (synchronizable)
threads.



GPU Architecture (Volta)

■ Processor Array

- executing many parallel threads efficiently by using many simple cores (CUDA cores)
- organizing streaming multiprocessors (SMs) each consisting of ...
 - 64 CUDA cores (INT+FP32)
 - 8 Tensor cores (4x4 matmul and acc)
 - Instruction and constant caches



High Performance Computing

Mapping Data to GPU Mem

```
int N = 1000;
```

```
int main()  
{
```

```
    int A = 0;
```

```
    int* B = (int*)malloc(sizeof(int)*N);
```

```
    #pragma omp target
```


```
    #pragma omp teams distribute parallel for
```

```
    for(int i = 0; i < N; i++){
```

```
        // A, B, N, and i exist here
```

```
        // B is a pointer – the data B point to DOES NOT exist here
```

```
    }  
}
```



Mapping Data to GPU Mem

```
int N = 1000;
```

```
int main()
```

```
{
```

```
    int A = 0;
```

```
    int* B = (int*)malloc(sizeof(int)*N);
```

```
    #pragma omp target map(tofrom:B[0:N])
```

```
    #pragma omp teams distribute parallel for
```


```
    for(int i = 0; i < N; i++){
```

```
        // A, B, N, and i exist here
```

```
        // B is a pointer – the data B point to DOES exist here
```

```
    }
```


```
}
```



Parallelization Scalability

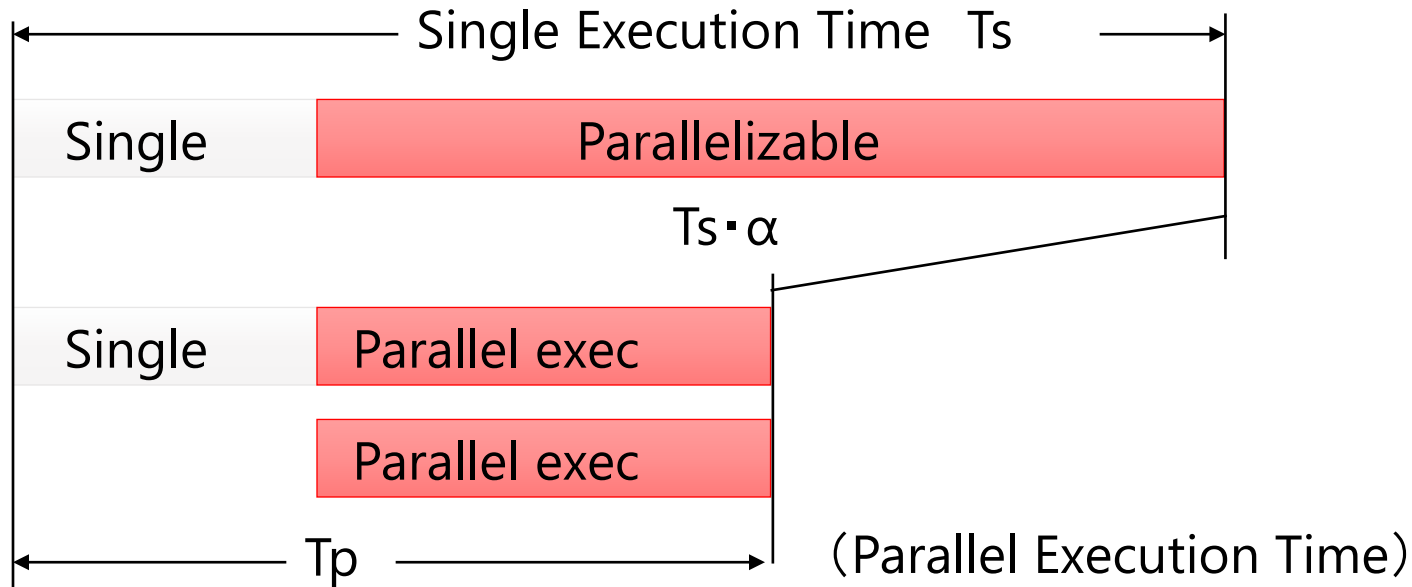
■ **Speedup ratio:** the ratio of parallel exec. to sequential exec.

- How many times does it become faster?



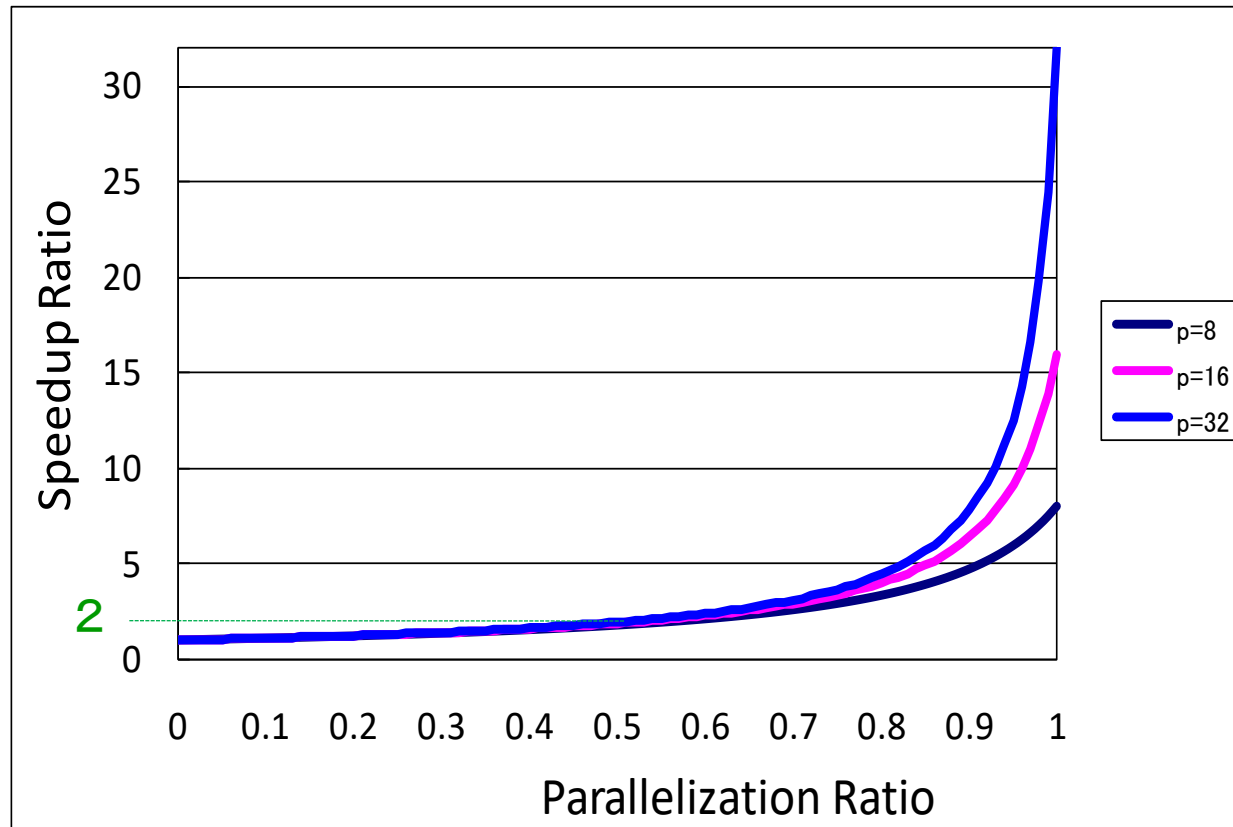
Parallelization Ratio

Parallelization Ratio α : The ratio of parallel exec.



$$\text{Speedup ratio} = \frac{T_s}{T_p} = \frac{1}{1 - \alpha + \alpha/n} \quad (\text{Amdahl's law})$$

Amdahl's Law



In the case of $\alpha=0.5$ (50%), the speedup ratio does not exceed 2.

= It is very important to select an algorithm with high parallelism.

What part should be accelerated?

- Pareto's laws (aka 80-20 rule)
- A small part of the code consumes a large part of the execution time.

■ Accelerate a time-consuming part!

- Significant reduction in the execution time.

98% of exec time consumed.

PROG.UNIT	FREQUENCY	EXCLUSIVE TIME[sec](%)	AVER.TIME [msec]	MOPS	MFLOPS	V.OP RATIO	AVER. V.LEN	VECTOR TIME	I-CACHE MISS	O-CACHE MISS	BANK CONF
calcPoisson()	20	0.909(98.2)	45.471	14384.2	5698.9	98.98	128.1	0.906	0.0000	0.0000	0.0022
calcBoundary_SqObject()	14578	0.012(1.3)	0.001	1333.1	218.0	94.84	22.7	0.009	0.0000	0.0000	0.0002
calcTantVelocity()	20	0.003(0.3)	0.128	23890.0	12568.5	99.68	118.0	0.003	0.0000	0.0000	0.0000
calcVelocity()	20	0.001(0.1)	0.044	13598.9	5773.6	99.06	118.0	0.001	0.0000	0.0000	0.0000
calcPoissonSourceTerm()	20	0.001(0.1)	0.032	14865.2	8021.5	98.38	118.0	0.000	0.0000	0.0001	0.0000
main	1	0.000(0.1)	0.490	225.9	0.1	0.00	0.0	0.000	0.0001	0.0000	0.0000

The performance info of a program compiled with `-pg` can be printed out by using the **gprof** command on a standard Linux system (e.g., AOBA-B).

Parallelization Overhead

Sequential



Parallel



time →

Parallelization Overhead

4 threads != 4 times faster

EX) thread creation and deletion



Parallelization Overhead

Sequential

CPU time

```
#pragma omp parallel for  
for (int i=0; ...) {  
#pragma omp parallel for  
for (int i=0; ...) {
```

Twice

Parallel

CPU time

CPU time

CPU time

CPU time

```
#pragma omp parallel  
{  
#pragma omp for  
for (int i=0; ...) {  
#pragma omp for  
for (int i=0; ...) {  
}
```

Once

Summary

■ Introduction to OpenMP Programming

- Data-parallel processing (work-sharing)
 - How to parallelize a for loop.
- Critical section
 - Shared variables and private variables.
- Task-parallel processing
 - Sections
 - Tasks
- Offloading
 - Target

■ Some keys to achieve better scalability

- Amdahl's law
- Pareto's law