



TOHOKU
UNIVERSITY



Cyberscience
Center

High Performance Computing

高性能計算論

Volume 6

Cyberscience Center, Tohoku Univ

Hiroyuki Takizawa

<takizawa@tohoku.ac.jp>

What you learnt so far (1/2)

■ Parallel Computers

- Shared-memory computers
- Distributed-memory computers
- Hierarchical (hybrid) systems
- Networks

What's Parallel Computer (1/3)

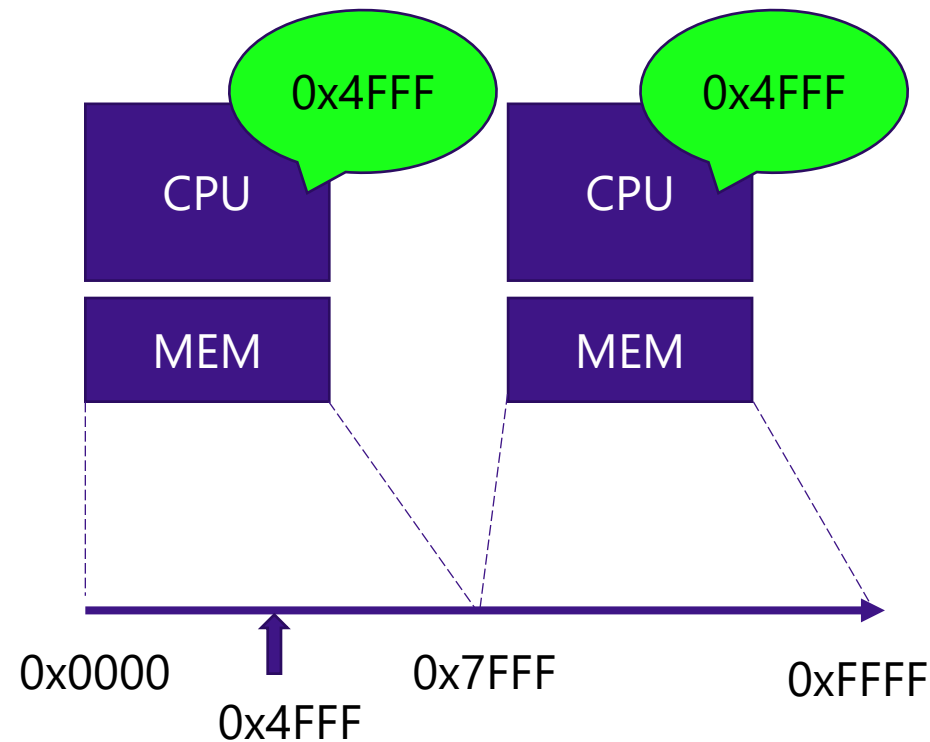
■ Parallel Computer

- A multi-processor computer system supporting parallel programming
- Two major categories of parallel computers
 - Distributed-memory parallel computers
 - Multiple computers and their interconnection network.
 - Employed to build a large-scale system
 - Shared-memory parallel computers
 - Symmetric multi-processor(SMP) and multicore/manycore.
 - Employed by most of current processors.

What's Parallel Computer (2/3)

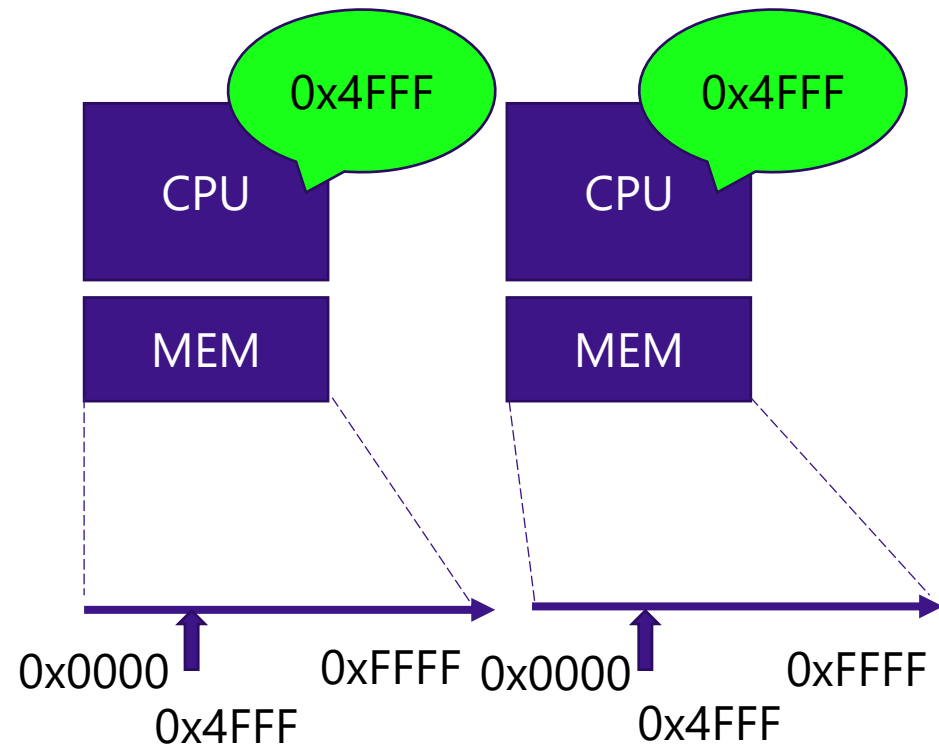
Shared-memory (NUMA)

Each memory device is mapped to a part of the memory space.



Distributed-memory

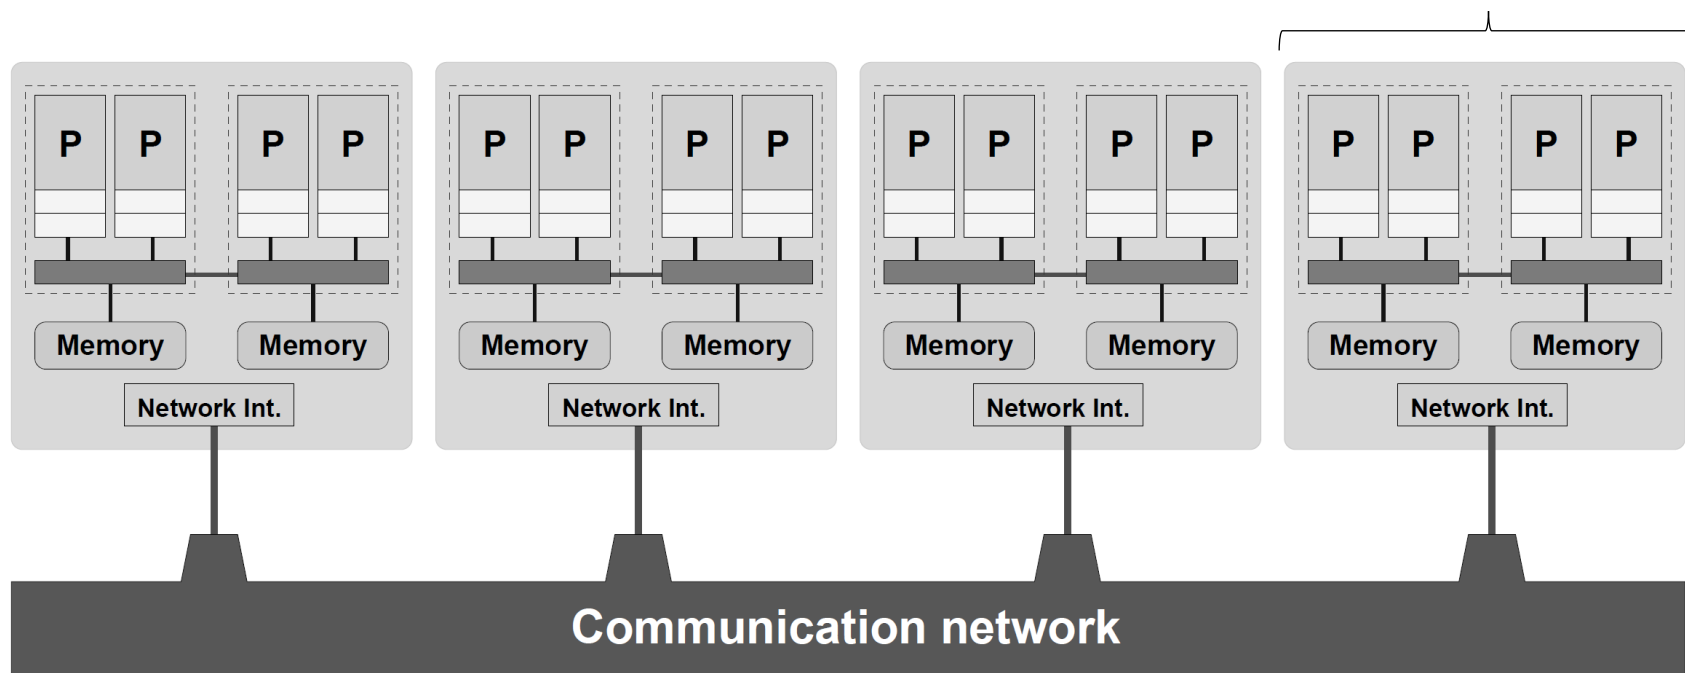
Each memory device has its own memory space.



What's Parallel Computer (3/3)

- Large-scale parallel computers
= mixture of shared and distrib.-parallel.

One OS instance manages a node.



In addition, each node may have accelerators such as GPUs.

What you learnt so far (2/2)

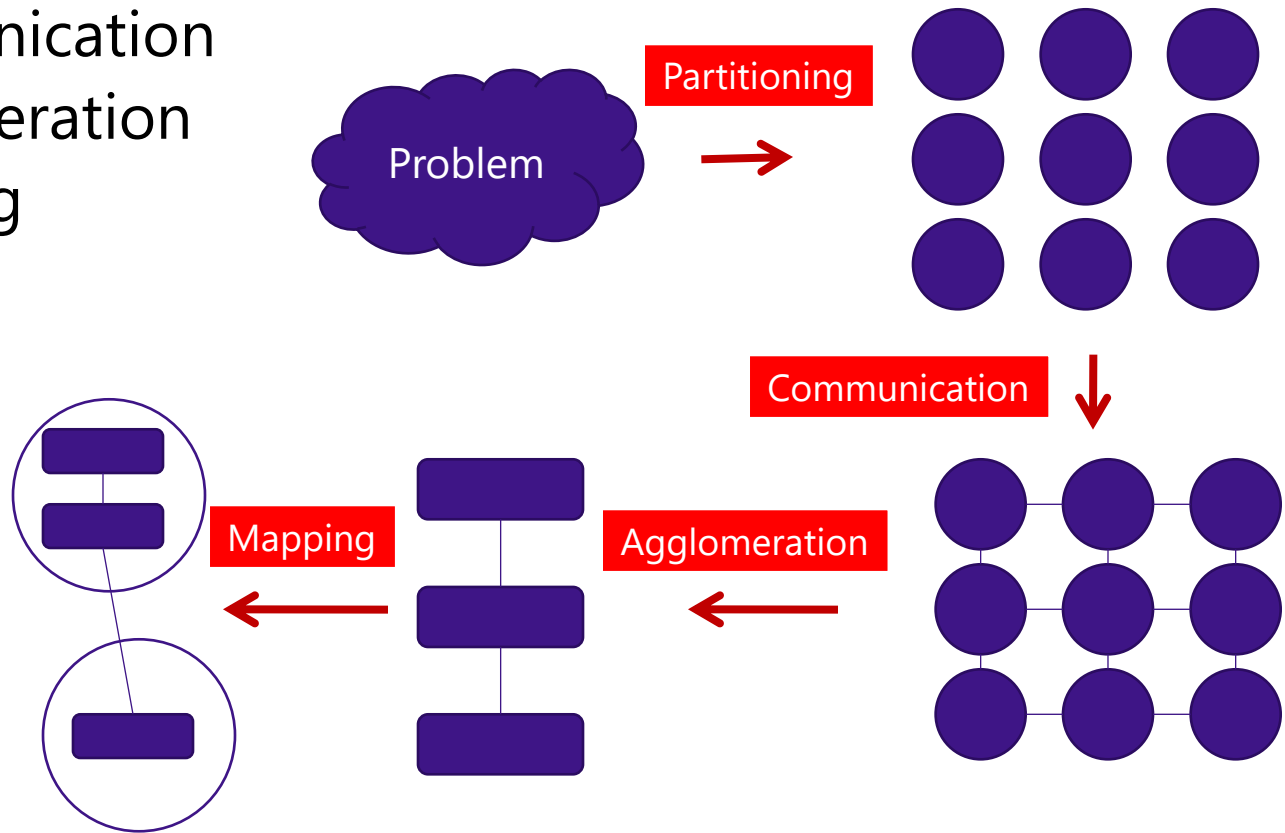
■ Parallel Algorithm Design

- Task/Channel Model
- Foster's Design Methodology
- Communication Patterns

Foster's Design Methodology

■ Four steps for designing parallel algorithms

- Partitioning
- Communication
- Agglomeration
- Mapping



What is MPI?

■ Message Passing Interface (**MPI**)

- Interface for parallel programs with message passing.
 - Multiple programs (**MPI processes**) run on a parallel computer.
 - Each MPI process has its own memory space.
 - MPI processes can pass their data to others if necessary.
 - MPI defines only the interface (not the implementation).
 - We do not need to care about how MPI processes actually communicate.

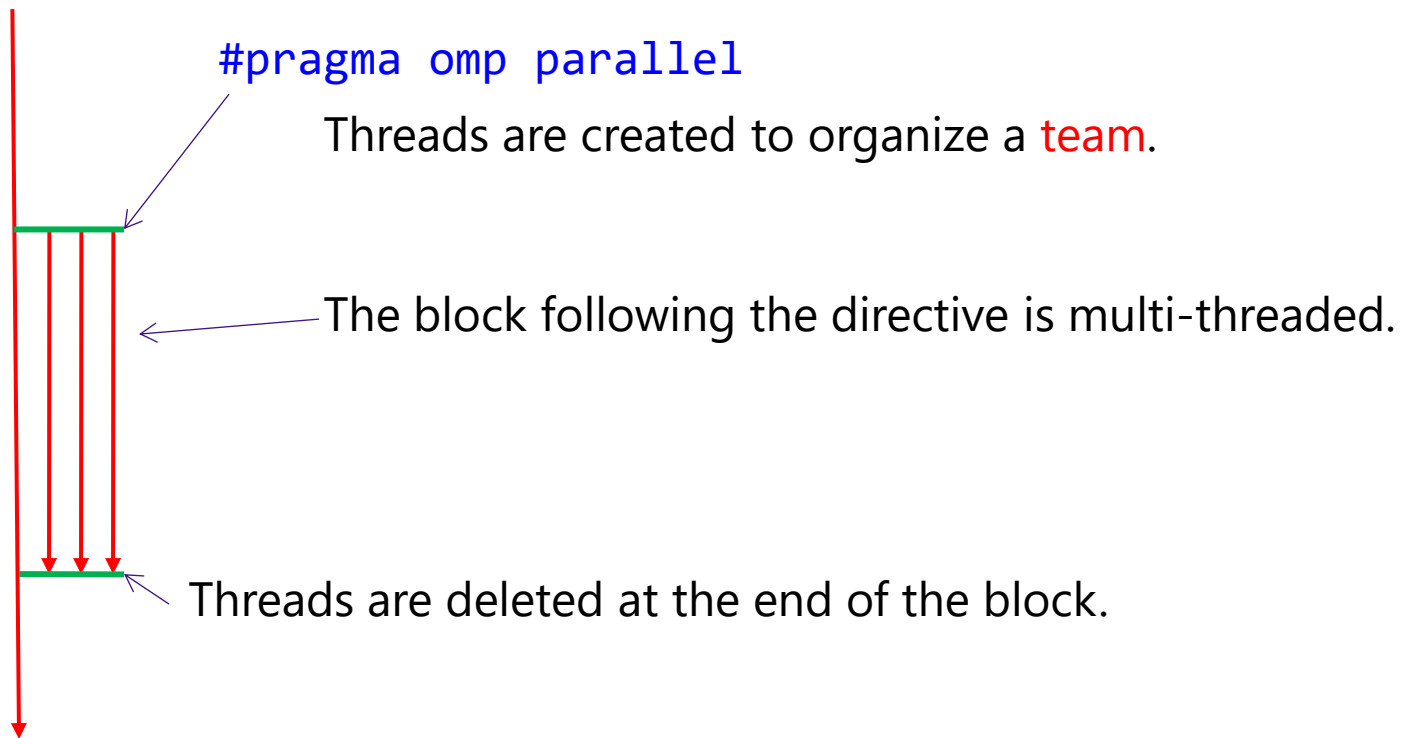
■ Major MPI Implementations

- MPICH (<http://www.mpich.org/>)
- Open MPI (<http://www.open-mpi.org/>)

What's OpenMP?

■ **Threads** are created/deleted on demand.

- Thread: an execution flow



Today's Topic


■ Advantages of using both MPI and OpenMP

- Case Study: Conjugate gradient method


■ Performance Prediction and Analysis

- Roofline model (memory bandwidth)
- Amdahl's Law (strong scaling ability)
- Gustafson-Barsis's Law (weak scaling ability)
- Karp-Flatt Metric (parallel performance evaluation)

C+MPI vs. C+MPI+OpenMP



C + MPI



C + MPI + OpenMP

C + MPI + OpenMP

- Lower communication overhead
- More portions of program may be practical to parallelize
- May allow more overlap of communications with computations

Case Study: Conjugate Gradient

- A is positive definite if for every nonzero vector x and its transpose x^T , the product $x^T A x > 0$
- If A is symmetric and positive definite, then the function

$$q(x) = \frac{1}{2} x^T A x - x^T b + c$$

has a unique minimizer that is solution to $Ax = b$

- Conjugate gradient is an iterative method that solves $Ax = b$ by minimizing $q(x)$

Case Study: Conjugate Gradient

■ Conjugate gradient method solves $Ax = b$

■ In our program we assume A is dense

- Matrix-vector multiplication
- Inner product (dot product)
 - Matrix-vector multiplication has higher time complexity


■ Methodology

- Start with a sequential program
- Profile functions to determine where most execution time spent
- Tackle most time-intensive function first

MPI Parallelization

A is decomposed and distributed over MPI processes

X is replicated for all MPI processes



Allgather communication is needed for all MPI processes to have the entire resultant vector (=new X).

Performance Profiling

- Compiler with `-pg` option
- Execute the program
- Use the `gprof` command to see the performance profile

```
$ mpicc -pg cg.c main.c MyMPI.c -o cg
$ mpirun -np 1 ./cg a-huge b-huge
$ ls
MyMPI.c a-huge b-huge cg cg.c gmon.out main.c
$ gprof ./cg | less

... Performance Profile Information ...
```


Result of Profiling MPI Program

Clearly our focus needs to be on function
matrix_vector_product

Flat profile:

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self s/call	total s/call	name
99.97	110.41	110.41	530	0.21	0.21	matrix_vector_product
0.05	110.46	0.05	1060	0.00	0.00	dot_product
0.02	110.48	0.02	10	0.00	11.05	cg
0.00	110.48	0.00	1063	0.00	0.00	my_malloc
0.00	110.48	0.00	530	0.00	0.00	create_mixed_xfer_arrays
0.00	110.48	0.00	530	0.00	0.00	replicate_block_vector
0.00	110.48	0.00	2	0.00	0.00	get_size
0.00	110.48	0.00	1	0.00	0.00	print_replicated_vector
0.00	110.48	0.00	1	0.00	0.00	print_subvector
0.00	110.48	0.00	1	0.00	0.00	read_replicated_vector
0.00	110.48	0.00	1	0.00	0.00	read_row_striped_matrix

Code for matrix_vector_product

Lines 90-103 in cg.c

```
void matrix_vector_product (int id, int p,
    int n, double **a, double *b, double *c)
{
    int    i, j;
    double tmp;          /* Accumulates sum */
    for (i=0; i<BLOCK_SIZE(id,p,n); i++) {
        tmp = 0.0;
        for (j = 0; j < n; j++)
            tmp += a[i][j] * b[j];
        piece[i] = tmp;
    }
    new_replicate_block_vector (id, p,
        piece, n, (void *) c, MPI_DOUBLE);
}
```

Adding OpenMP directives

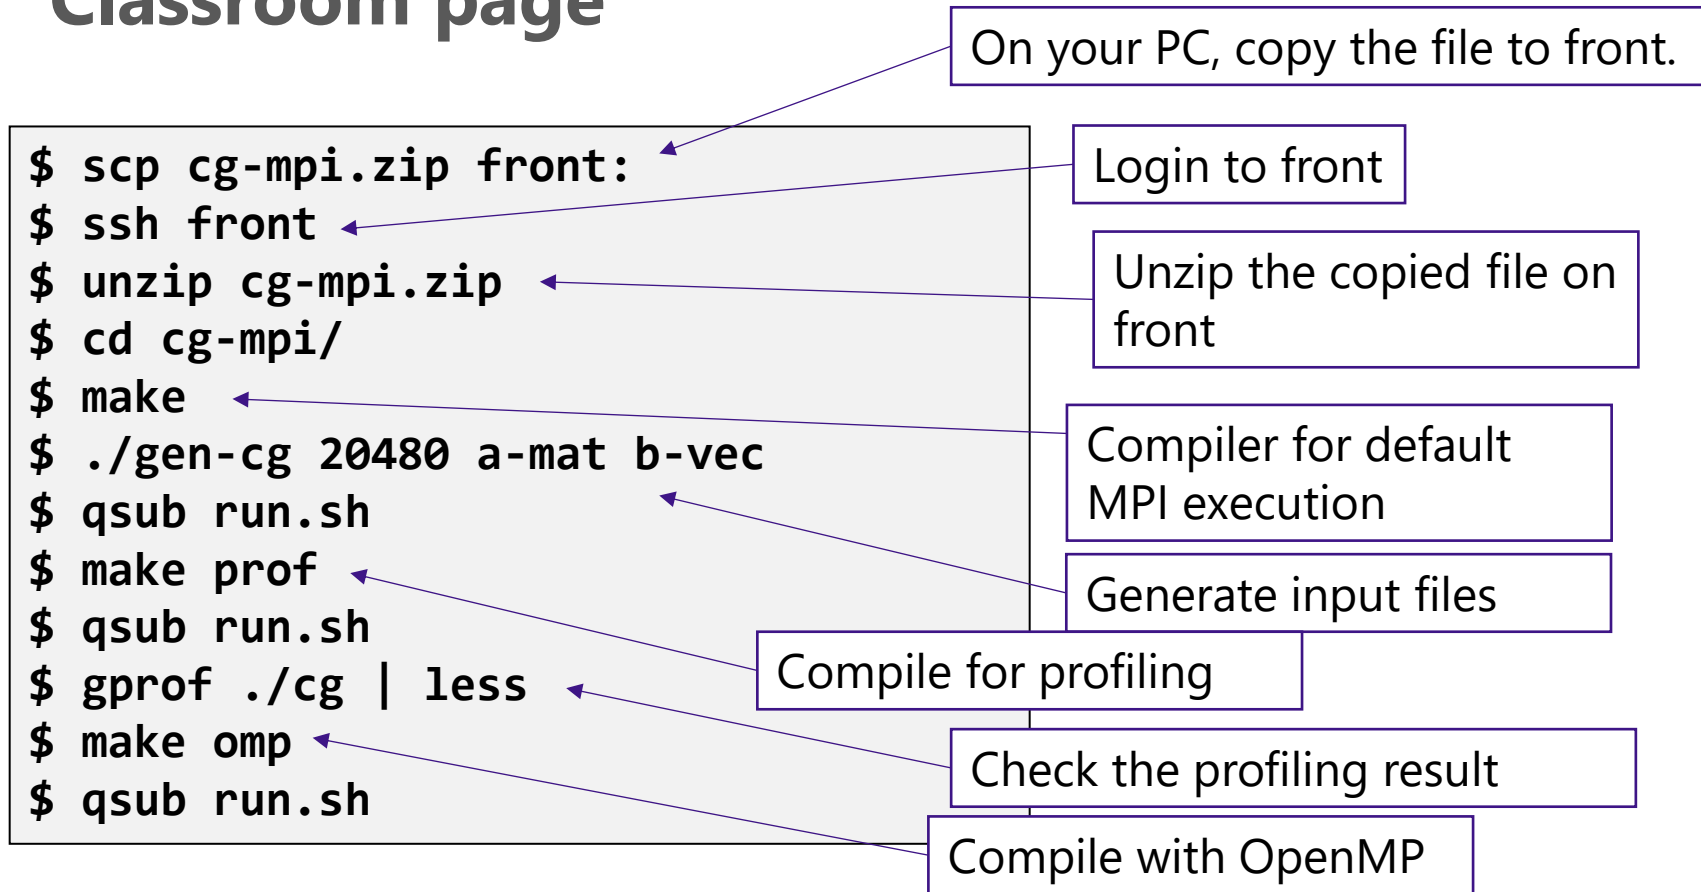
- Want to minimize fork/join overhead by paralleling the outermost possible loop
- Outer loop may be executed in parallel if each thread has a private copy of tmp and j

```
#pragma omp parallel for private(j,tmp)  
    for (i=0; i<BLOCK_SIZE(id,p,n); i++) {
```

We transformed a C+MPI program to a C+MPI+OpenMP program by **adding only one line** to our program!

Excercise

■ Download the zip file from the Google Classroom page



You can change the number of threads, and the number of processes. See run.sh.

Today's Topic

■ Advantages of using both MPI and OpenMP

- Case Study: Conjugate gradient method

■ Performance Prediction and Analysis

- Roofline model (memory bandwidth)
- Amdahl's Law (strong scaling ability)
- Gustafson-Barsis's Law (weak scaling ability)
- Karp-Flatt Metric (parallel performance evaluation)

Performance Prediction and Analysis

■ Why do we need **performance prediction** of a parallel algorithm?

- Because it helps us confirm whether an expected performance has been achieved or not.

■ The purposes of **performance analysis**

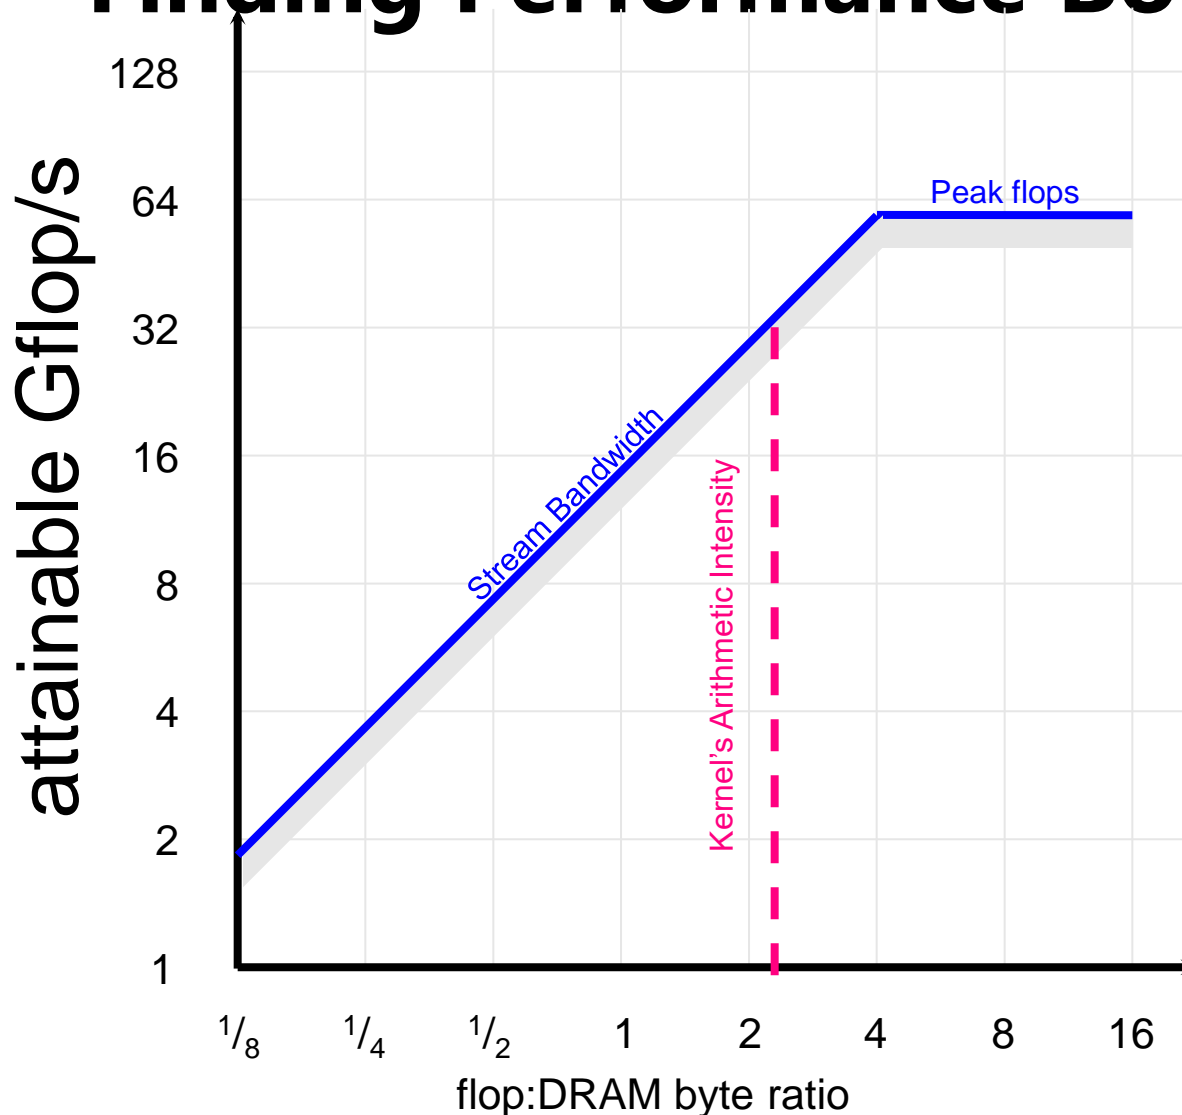
- To understand the barriers to high performance
- To predict how much improvement can be achieved by increasing the number of processors/cores.

Finding Performance Bottleneck (1/2)

■ Roofline model (Williams, 2009)

- Suppose a computer of **B** [bytes/s] and **F** [flop/s]. Then, **what is the expected flop/s rate on simple vector addition?**
 - Assume each vector element is a 64-bit floating point value.
 - Two operands are needed for one addition, and one result is produced (Assume each is a 64-bit floating point value).
→ 24 bytes for one addition (2 inputs and 1 output).
- = **Arithmetic intensity** of vector addition is $1/24$
- The computer can achieve its peak flop/s rate, F , only if B is 24 times larger than F . Otherwise, the memory bandwidth limits the theoretically achievable flop/s rate (MB = performance bottleneck).

Finding Performance Bottleneck (2/2)



Roofline model (Williams, 2008)

- ❖ Machines have finite memory bandwidth
- ❖ Apply a Bound and Bottleneck Analysis
- ❖ Still Unrealistically optimistic model

$$\text{Gflop/s(AI)} = \min \begin{cases} \text{Peak Gflop/s} \\ \text{StreamBW} * \text{AI} \end{cases}$$

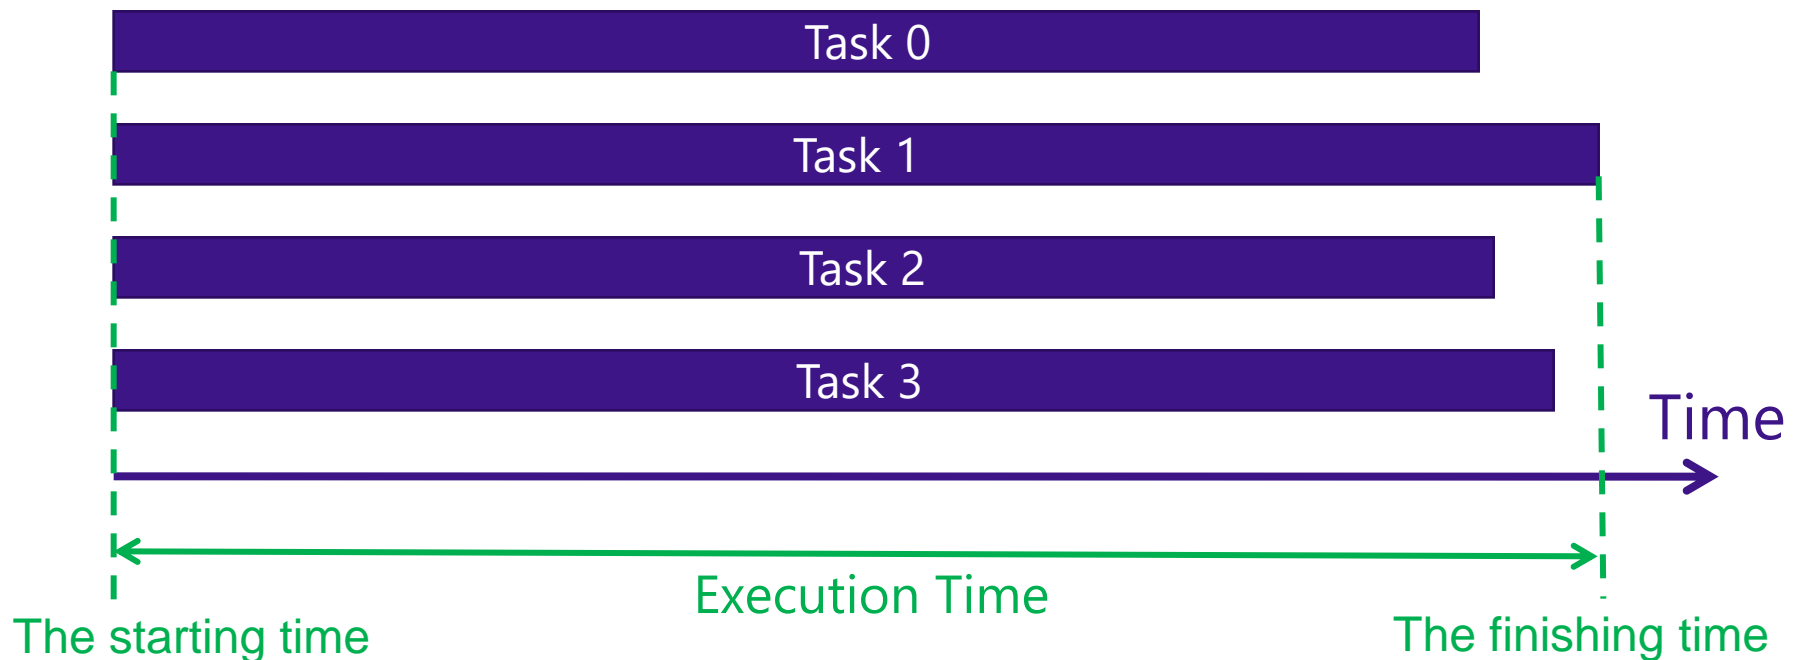
Figure by William 2008

Execution Time

■ Execution time of a parallel algorithm:

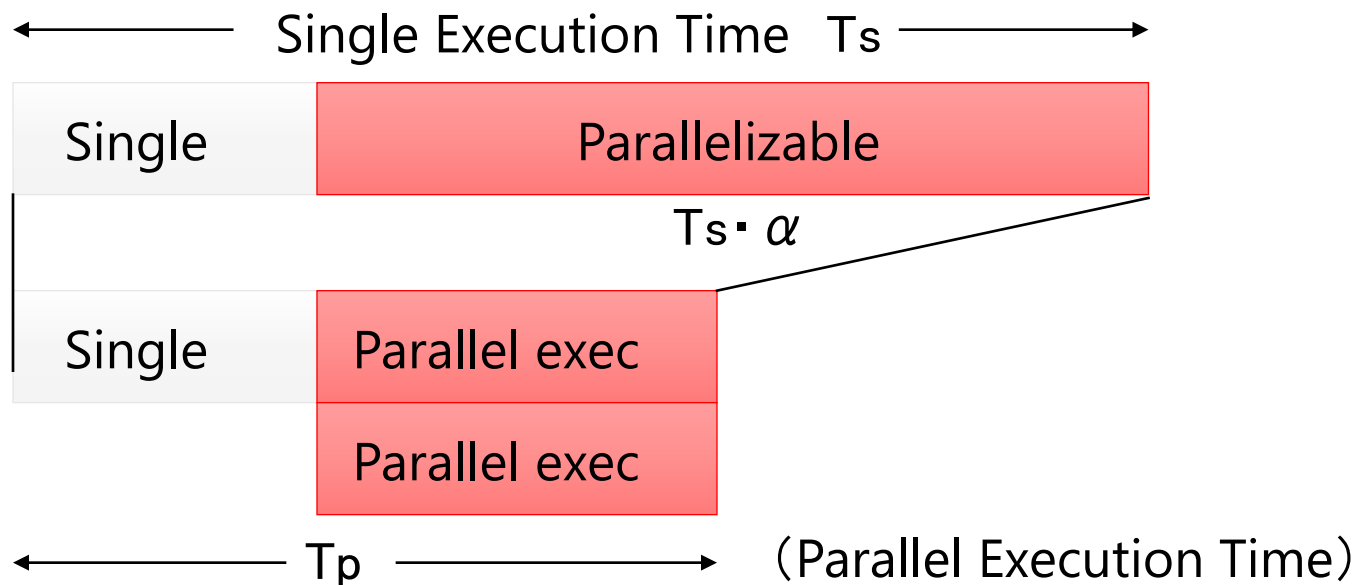
- The period of time during any task is active.

(Shorter is Better.)



Parallelization Ratio

Parallelization Ratio α : The ratio of parallel exec.



$$\text{Speedup ratio} = \frac{T_s}{T_p} = \frac{1}{1 - \alpha + \alpha/n} \quad (\text{Amdahl's law})$$

Parallelization Overhead

Sequential



Parallel



Parallelization Overhead
4 threads != 4 times faster

EX) thread creation and deletion

time →

Speedup and Efficiency

■ Speedup

- The ratio between sequential execution time and parallel execution time

$$\text{Speedup} = (\text{Sequential Exec. Time}) / (\text{Parallel Exec. Time})$$

■ Efficiency

- A measure of processor utilization

$$\text{Efficiency} = \text{Speedup} / (\text{Number of Processors})$$

■ Scalability

- A measure of parallel system's ability to increase the performance as the number of processor increases.

Why Ideal Speedup so difficult?

■ linear speedup is prevented by ...

- Serial operations
- Communication operations
- Process start-up
- Imbalanced workloads
- Architectural limitations

Amdahl's Law

■ Operations performed by parallel algorithm

- Computations that must be performed sequentially
- Computations that can be performed in parallel
- Parallelization overhead
 - Communication operations and redundant computations

■ Upper bound of speedup ratio for a fixed problem size (Amdahl's law)

$$\Psi(n, p) \leq \frac{\sigma(n) + \varphi(n)}{\sigma(n) + \varphi(n)/p + \kappa(n, p)} \leq \frac{\sigma(n) + \varphi(n)}{\sigma(n) + \varphi(n)/p} = \frac{1}{f + (1-f)/p}$$

$\Psi(n, p)$: speedup for a problem of size n on p processors

$\sigma(n)$: sequential portion of the computation

$\varphi(n)$: parallel portion of the computation

$\kappa(n, p)$: parallelization overhead

p : the number of processors

n : the problem size

f : the fraction of sequential computation = $\sigma(n)/(\sigma(n) + \varphi(n))$

Example 1

- 95% of a program's execution time occurs inside a loop that can be executed in parallel. What is the maximum speedup we should expect from a parallel version of the program executing on 8 CPUs?

Example 2

- 20% of a program's execution time is spent within inherently sequential code. What is the limit to the speedup achievable by a parallel version of the program?

Considering the overhead

■ Limitation of Amdahl's law


- The parallelization overhead such as communication time is ignored.
- Typically, the overhead is lower complexity than parallel portion of the computation
 - The time of parallel computation grows faster than the overhead as the problem size increases.

■ EX) Suppose a program

- sequential comp. : $O(n)$ e.g. $18000+n$
- parallel comp. : $O(n^2)$ e.g. $(n^2/100)/p$
- comm. : $O(n + \log p)$
e.g. $10000 \text{ ceil}(\log p) + (n/10)$

Illustration of Amdahl Effect

Speedup is usually an increasing function of the problem size.



Amdahl Effect

- Typically $\kappa(n,p)$ has lower complexity than $\phi(n)/p$
- As n increases, $\phi(n)/p$ dominates $\kappa(n,p)$
- As n increases, speedup increases

Review of Amdahl's Law

- Treats problem size as a constant
- Shows how execution time decreases as number of processors increases

The scalability with a fixed problem size is so-called **strong scalability**

Gustafson-Basis's Law

- What happens if the time is limited and the problem size increases with the number of processors?
 - Amdahl's effect: sequential fraction of a computation typically decreases as the problem size increases.
 - Increasing the number of processors enables us to increase the problem size solved in the time.
- Solving a problem of size n **using p processors** (=parallel execution), and s is the fraction of execution time spent in serial part.
→ $s = \sigma(n)/(\sigma(n)+\phi(n))$

$$\Psi(n, p) \leq \frac{\sigma(n) + p\phi(n)}{\sigma(n) + \frac{p\phi(n)}{p}} = p + (1 - p) \frac{\sigma(n)}{\sigma(n) + \phi(n)} = p + (1 - p)s$$

"Scaled speedup" improves with p (Gustafson-Basis's Law) = generally, too optimistic!

Karp-Flatt Metric

■ Experimentally determined serial fraction e

$T(n, p) = \sigma(n) + \varphi(n) / p + \kappa(n, p)$ \longleftarrow Total parallel execution time

$T(n, 1) = \sigma(n) + \varphi(n) + \kappa(n, 1)$ \longleftarrow Total serial execution time

Definition: experimentally determined serial fraction

$$e = (\sigma(n) + \kappa(n, p)) / T(n, 1)$$

$$\Rightarrow \sigma(n) + \kappa(n, p) = T(n, 1)e$$

$$T(n, p) = T(n, 1)e + T(n, 1)(1 - e) / p$$

$$T(n, 1) = T(n, p)\psi(n, p)$$

$$\Rightarrow T(n, p) = T(n, p)\psi e + T(n, p)\psi(1 - e) / p$$

$$\Rightarrow 1 = \psi e + \psi(1 - e) / p$$

$$\Rightarrow e = \frac{1/\psi - 1/p}{1 - 1/p} \quad (\text{smaller } e \text{ means better parallelization})$$

Karp-Flatt Metric (Cont'd)

Experimentally Determined Serial Fraction

$$e = \frac{\sigma(n) + \kappa(n, p)}{\sigma(n) + \varphi(n)}$$

Inherently serial component
of parallel computation +
processor communication and
synchronization overhead

$$e = \frac{1/\psi - 1/p}{1 - 1/p}$$

Single processor execution time

Karp-Flatt Metric (Cont'd)

■ Experimentally determined serial fraction is useful because ...

- Takes into account parallelization overhead
- Detects other sources of overhead or inefficiency ignored in speedup model
 - Process startup time
 - Process synchronization time
 - Imbalanced workload
 - Architectural overhead

Example 1

p	2	3	4	5	6	7	8
Ψ	1.8	2.5	3.1	3.6	4.0	4.4	4.7

What is the primary reason for speedup of only 4.7 on 8 CPUs?

e	0.1	0.1	0.1	0.1	0.1	0.1	0.1
---	-----	-----	-----	-----	-----	-----	-----

Since e is constant, large serial fraction is the primary reason.

Example 2

p	2	3	4	5	6	7	8
Ψ	1.9	2.6	3.2	3.7	4.1	4.5	4.7

What is the primary reason for speedup of only 4.7 on 8 CPUs?

e	0.070	0.075	0.080	0.085	0.090	0.095	0.100
---	-------	-------	-------	-------	-------	-------	-------

Since e is steadily increasing, overhead is the primary reason.

Today's Topic

■ Advantages of using both MPI and OpenMP

- Case Study: Conjugate gradient method

■ Performance Prediction and Analysis

- Roofline model (memory bandwidth)
- Amdahl's Law (strong scaling ability)
- Gustafson-Barsis's Law (weak scaling ability)
- Karp-Flatt Metric (parallel performance evaluation)